

ISOVALENT

Zero Trust Networking in Kubernetes

Using Cilium, Hubble & Grafana

Raphaël Pinson | @raphink | @raphink@mastodon.social
Solutions Architect, Isovalent





Key Questions when Enterprises Adopting Network Policy

- Where to start?
- How to troubleshoot Network Policies?
- How to stay up-to-date with Network Policies while applications evolve?
- How to prevent application teams from simple allowing everything?
- How can security teams prove that default deny policies are enforced?
- How can security teams be alerted on denied connections?

There is no “Easy Road” for Adopting Network Policy

HIGH RISK

- **Default Deny**
 - Each service-to-service communication explicitly allowed
 - But high chance of misconfiguration ⇒ application **unavailability** and **high adoption friction**
- Better Approach: Focus on **Risk Reduction**
 - Define **metrics** for Risk Exposure
 - Focus on the **most security sensitive** namespaces first
 - Leverage **network observability** to identify which network policy patterns easily reduce risk with minimal friction.

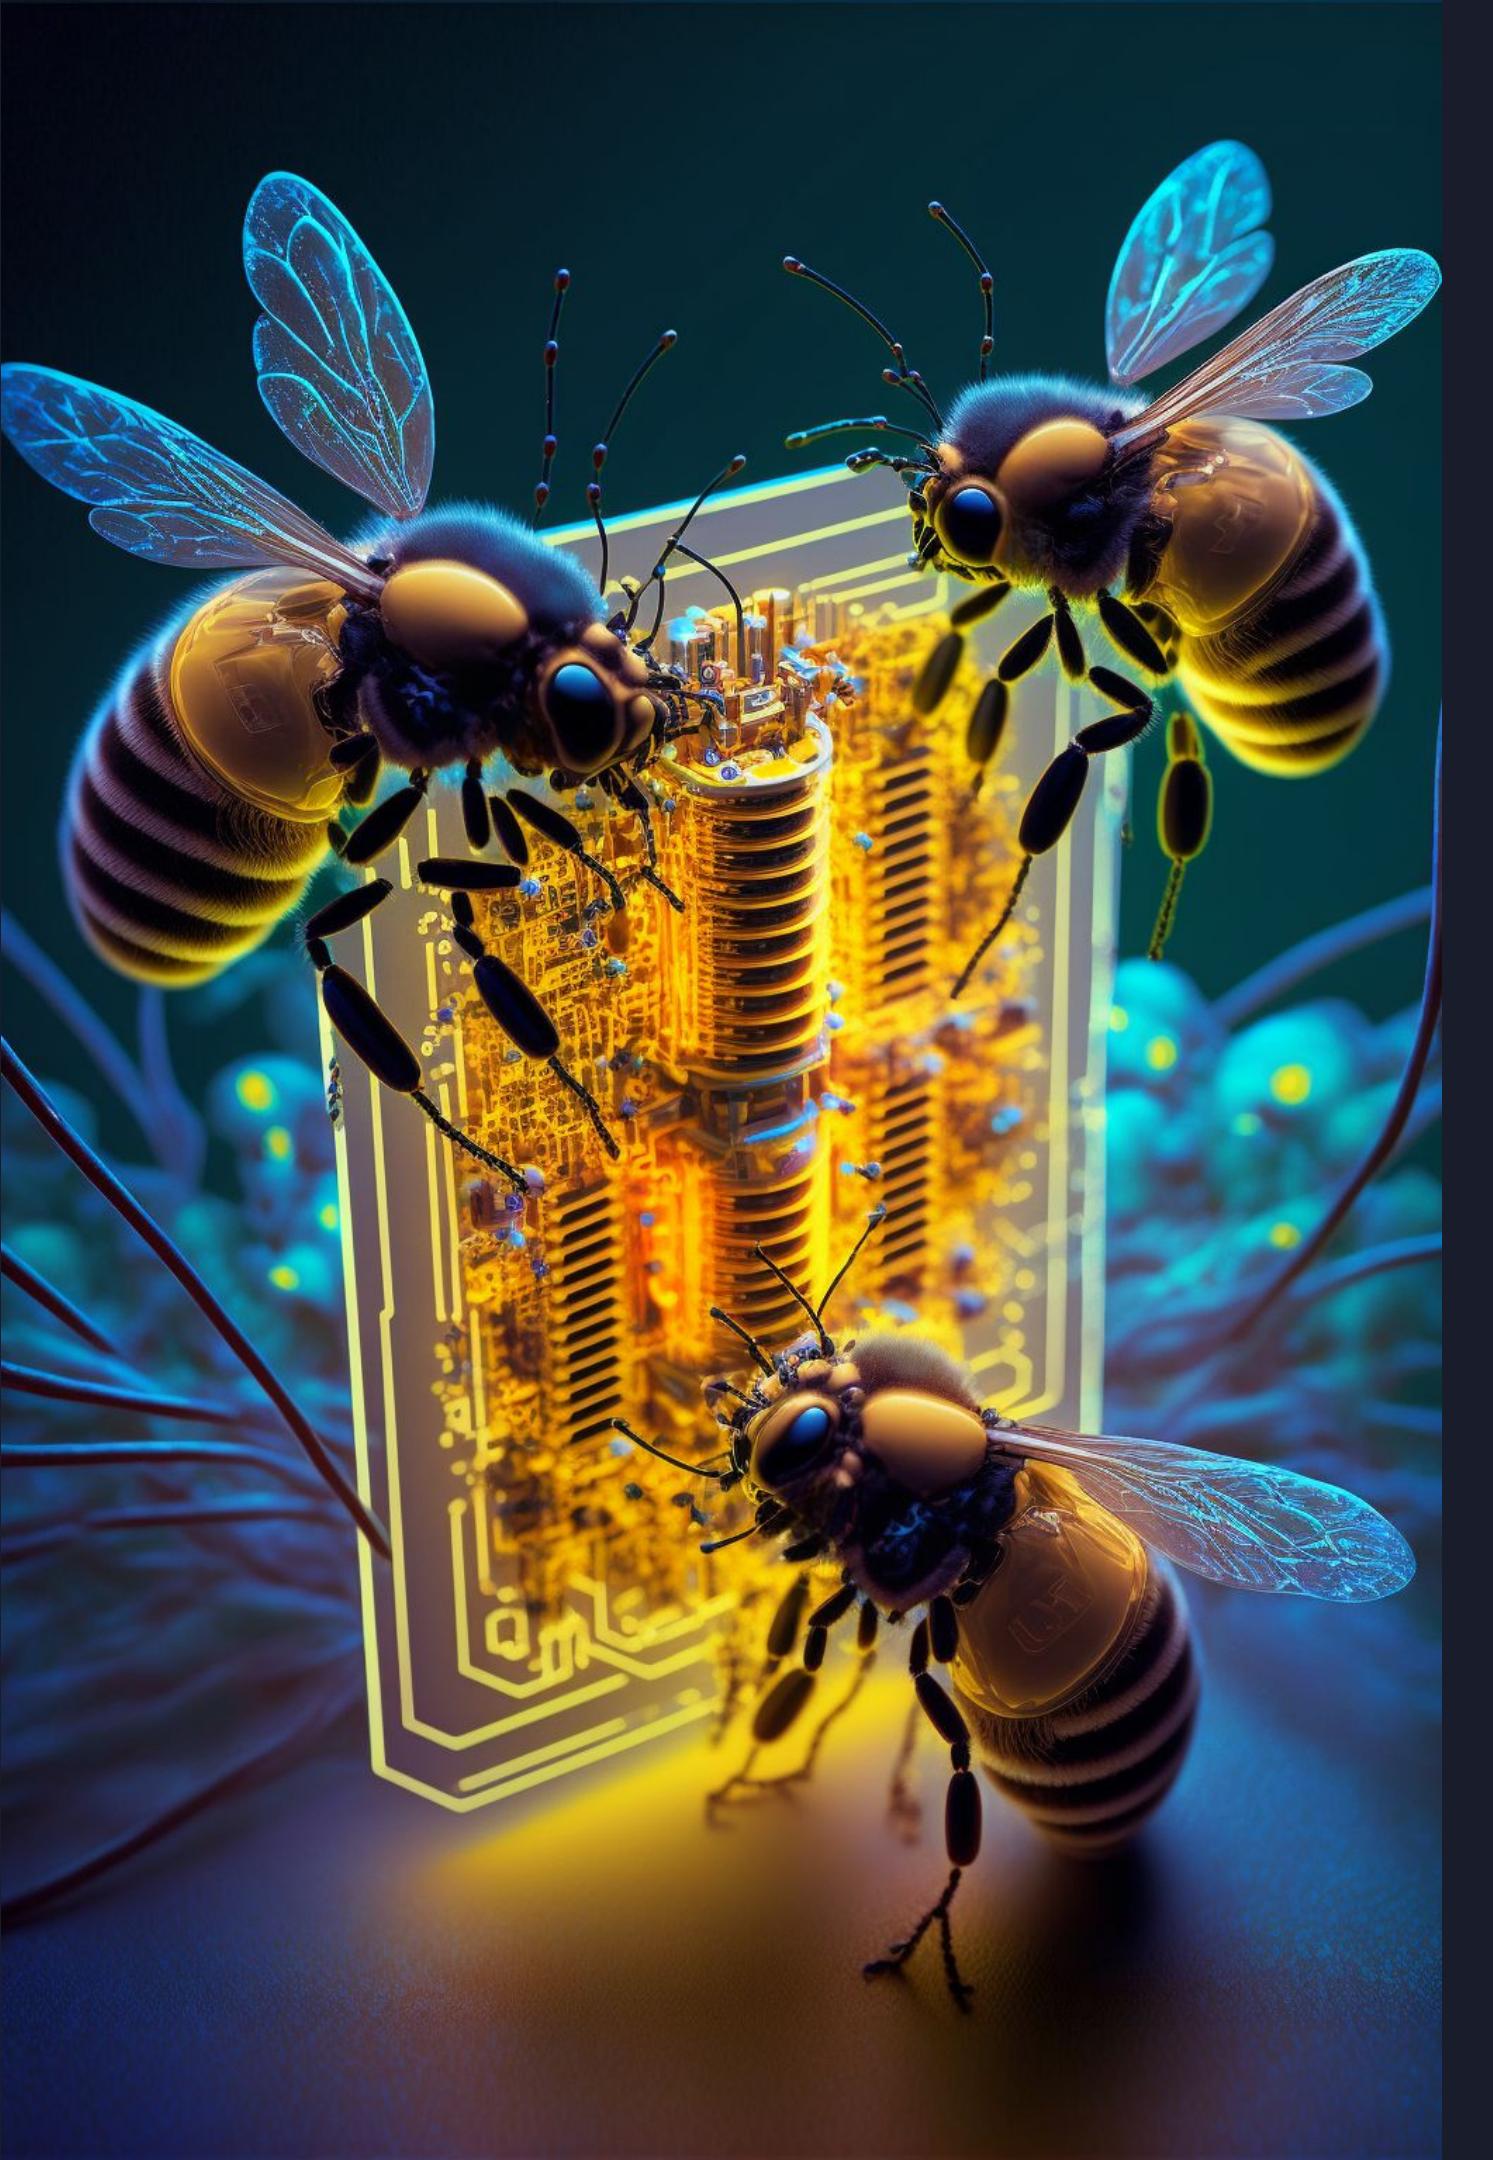


Who am I

Raphaël Pinson

Solutions Architect @ Isovalent
CNCF Ambassador





Zero Trust Networking

Using Cilium, Hubble & Grafana

- Cilium & eBPF
- Network Policies Strategies
- Network Policies Observability

What this talk is NOT about: Network Policy 101



ISOVALENT

Labs

isovalent.com/labs

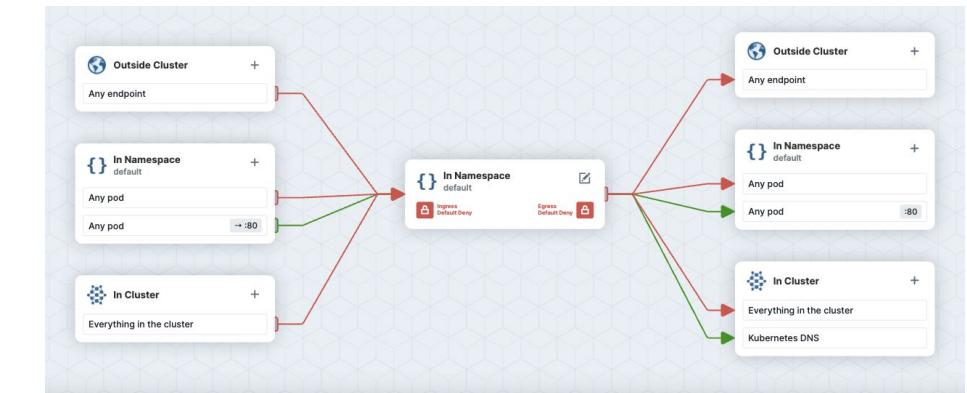


Tutorials

Getting started guides and examples.

cilium.io

networkpolicy.io



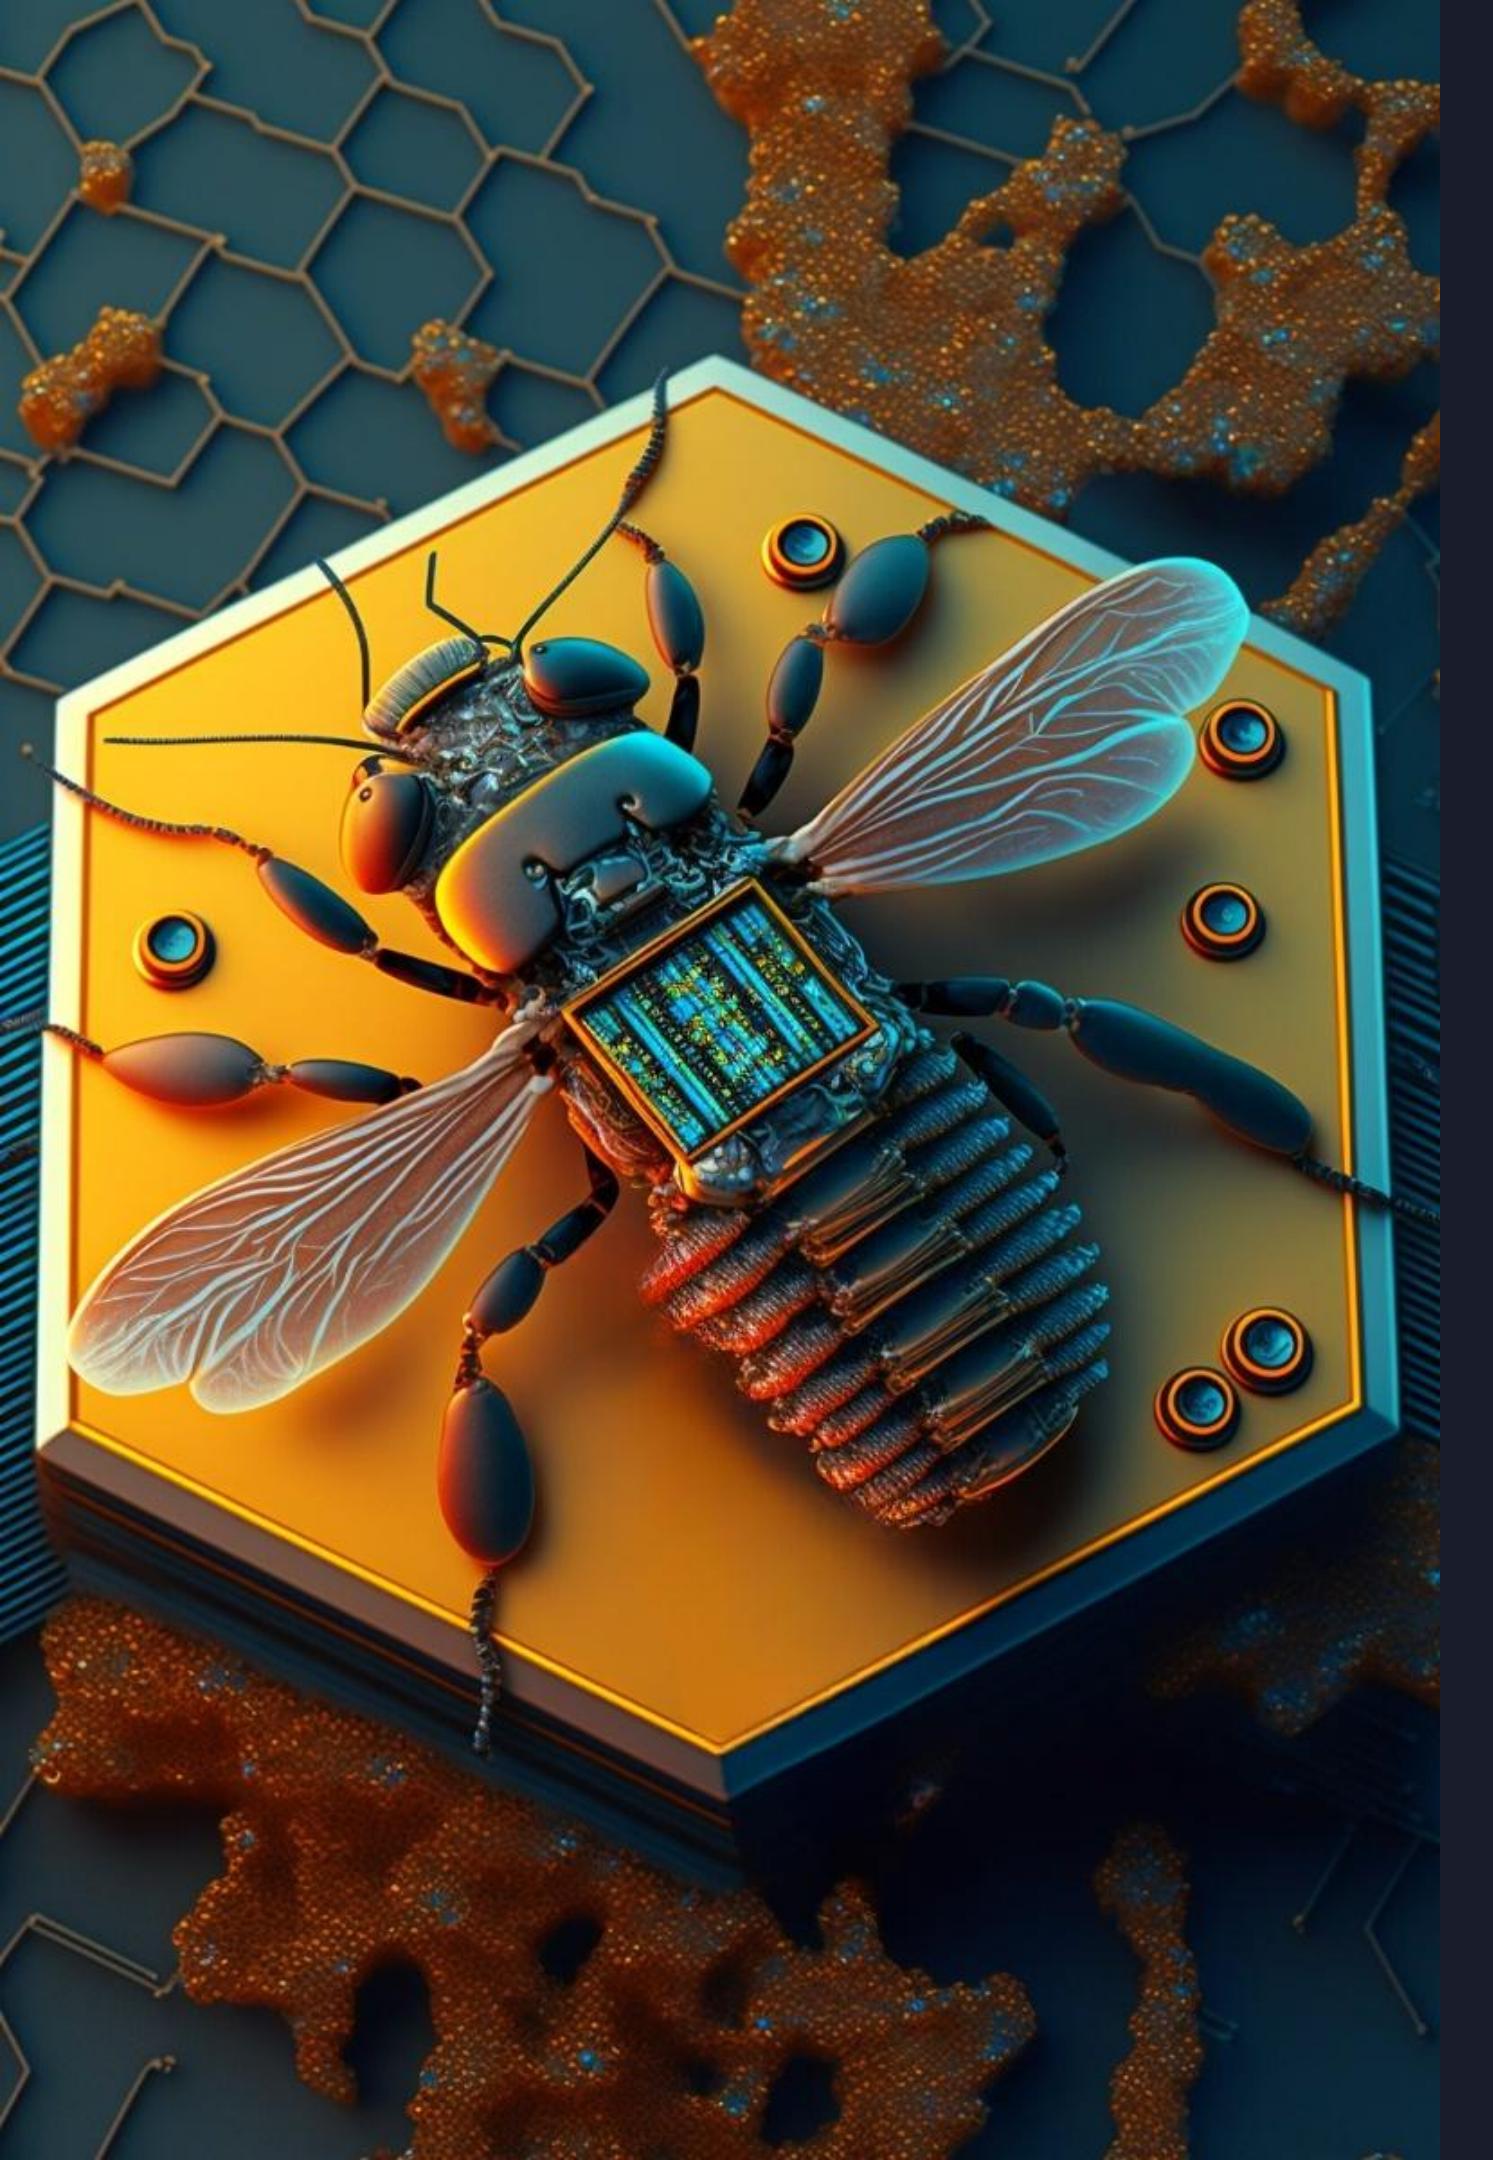
Tools

Learn how to create Network Policies for Kubernetes

editor.cilium.io

This session IS about strategies to overcome the operational & organizational challenges of adopting Network Policy in complex enterprise environments.

ISOVALENT



Zero Trust Networking

Using Cilium, Hubble & Grafana

- Cilium & eBPF
- Network Policies Strategies
- Network Policies Observability



- Open Source Projects

I SOVALENT

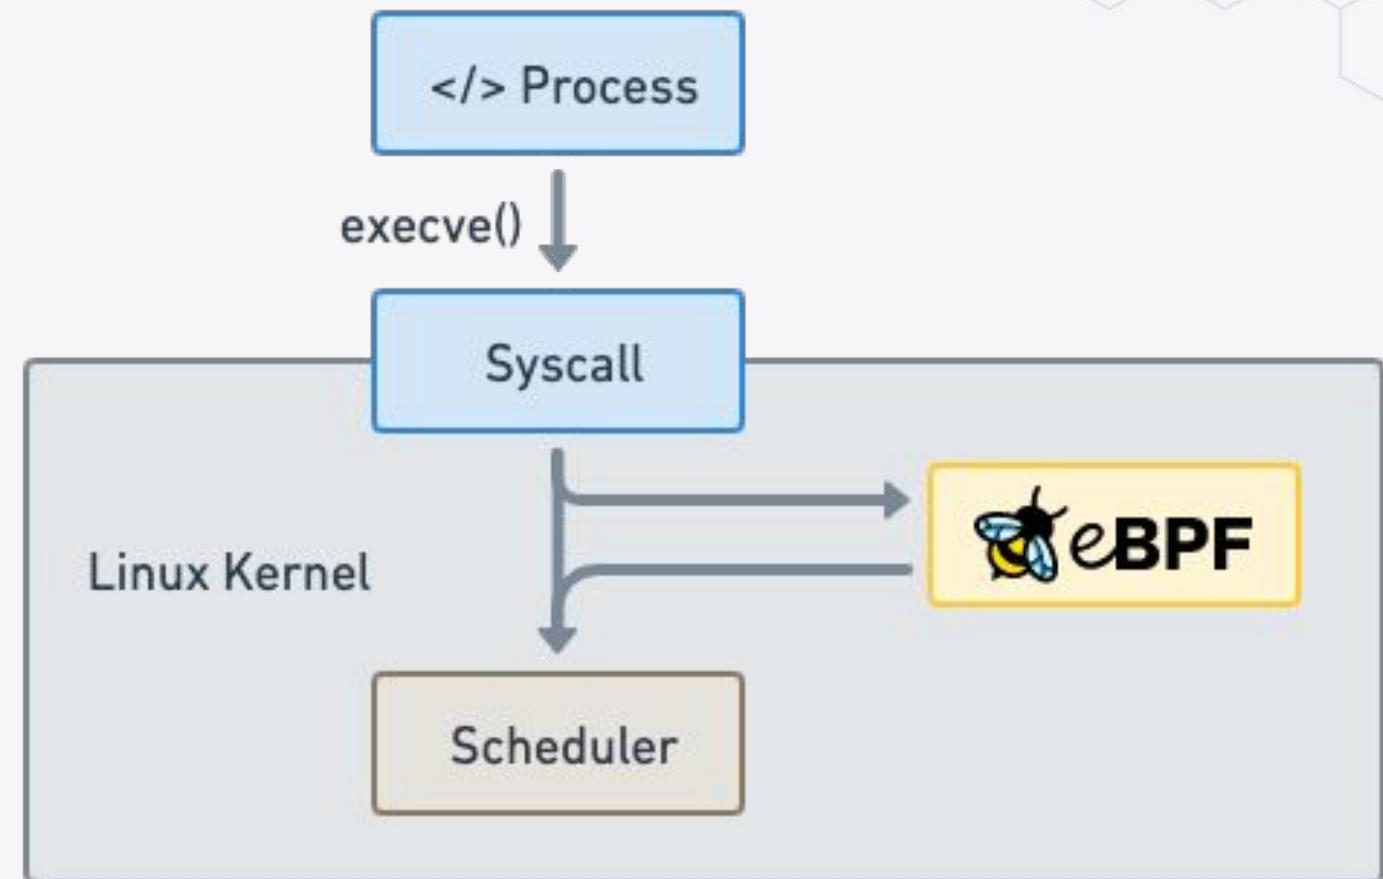
- Company behind Cilium
- Provides Cilium Enterprise





Makes the Linux kernel
programmable in a
secure and efficient way.

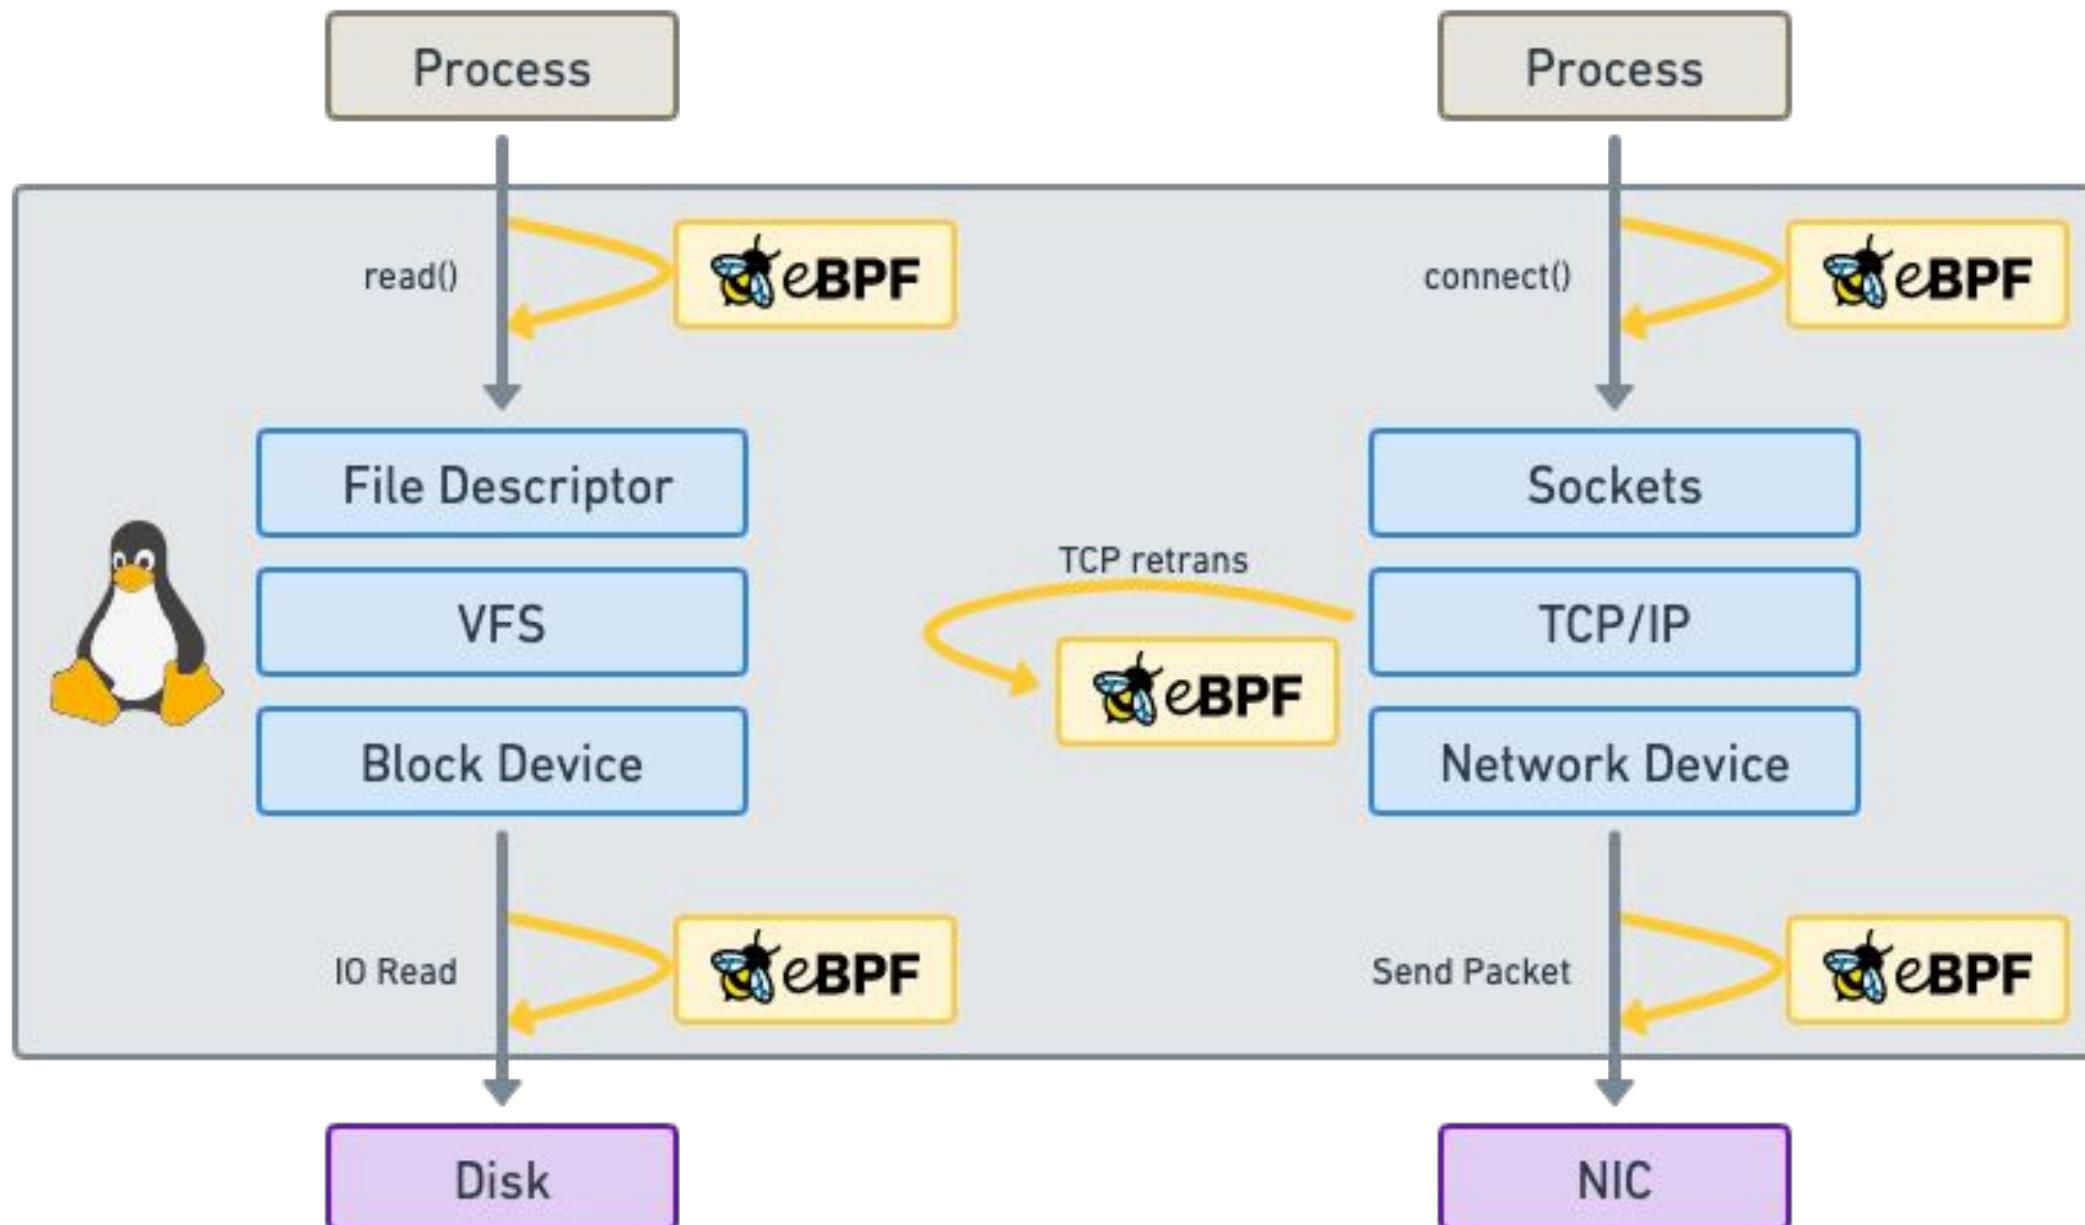
*“What JavaScript is to the
browser, eBPF is to the
Linux Kernel”*



```
int syscall__ret_execve(struct pt_regs *ctx)
{
    struct comm_event event = {
        .pid = bpf_get_current_pid_tgid() >> 32,
        .type = TYPE_RETURN,
    };
    bpf_get_current_comm(&event.comm, sizeof(event.comm));
    comm_events.perf_submit(ctx, &event, sizeof(event));

    return 0;
}
```

Run eBPF programs on events



Attachment points

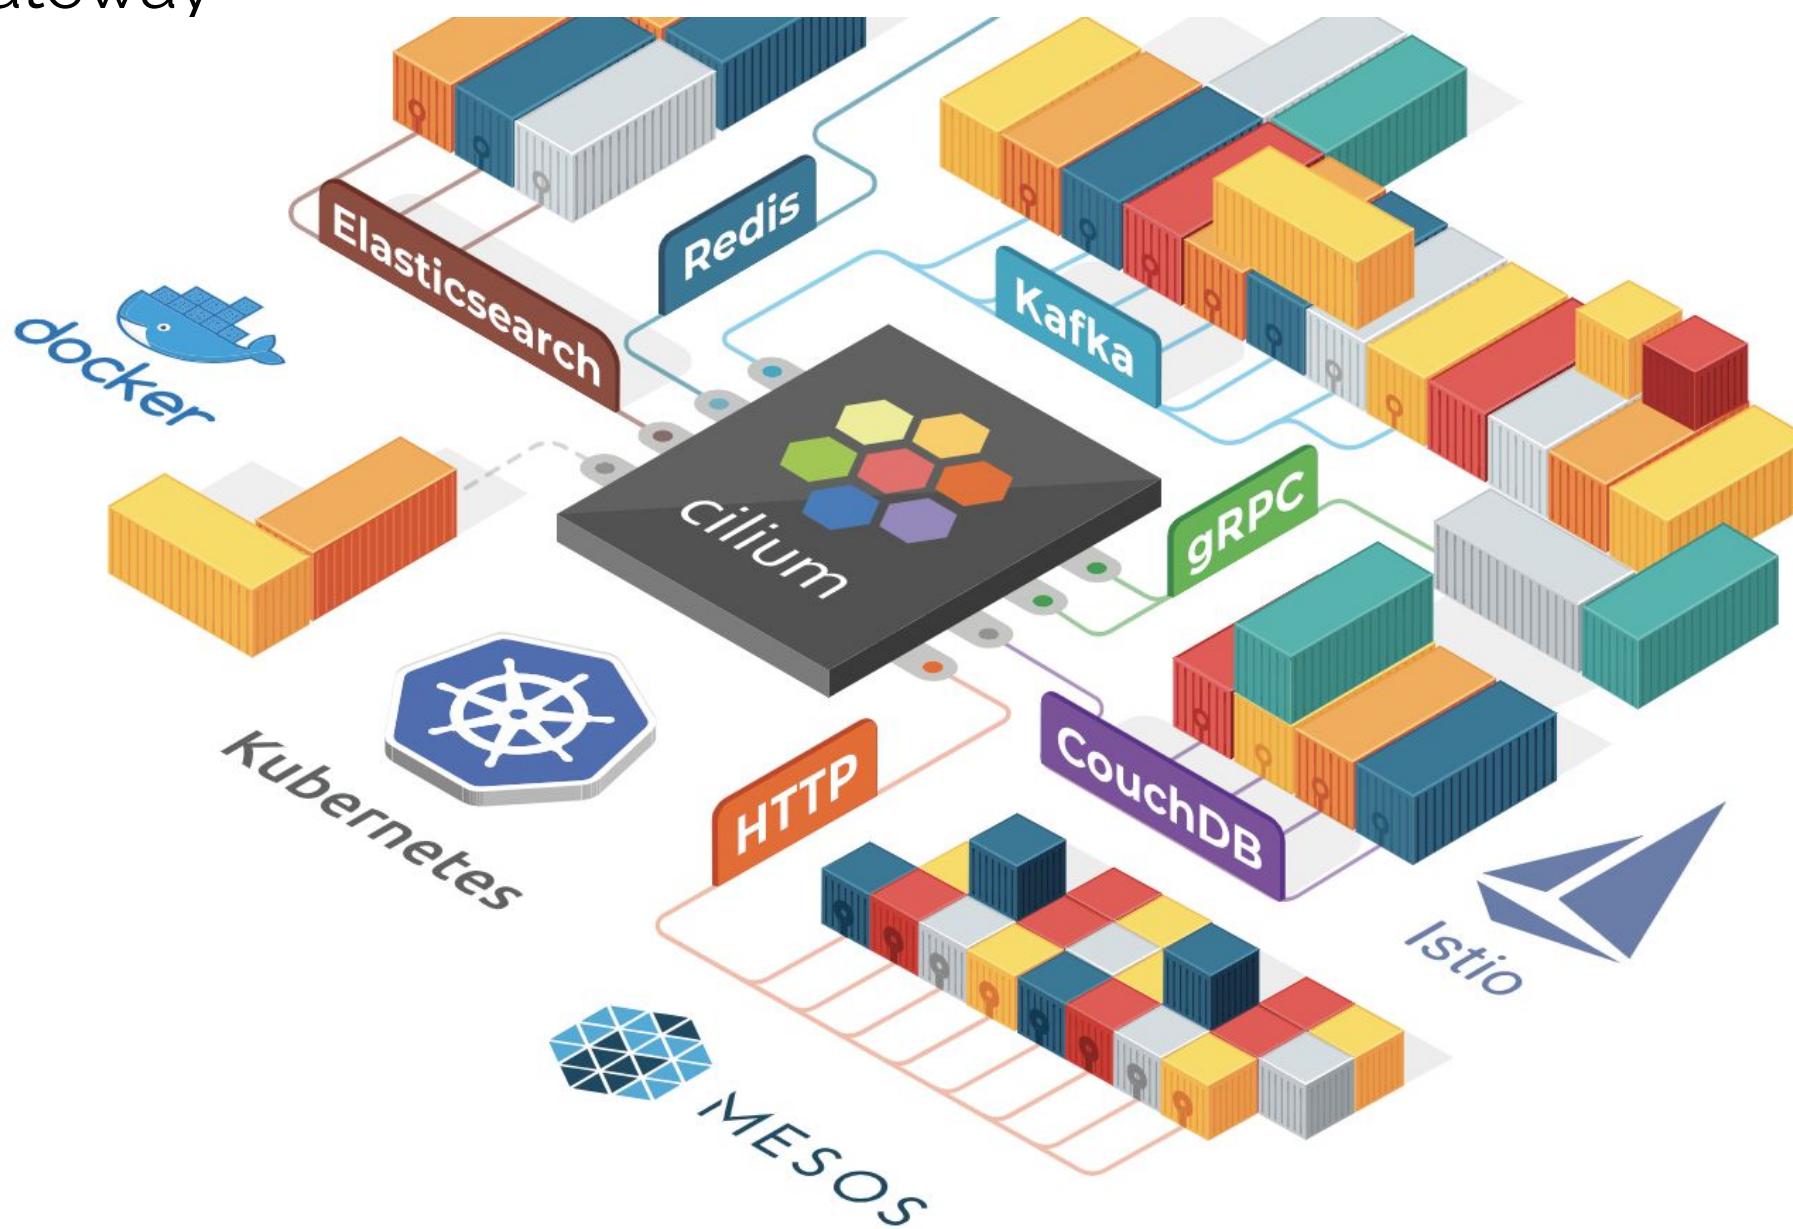
- Kernel functions (kprobes)
- Userspace functions (uprobe)
- System calls
- Tracepoints
- Sockets (data level)
- Network devices (packet level)
- Network device (DMA level) [XDP]
- ...

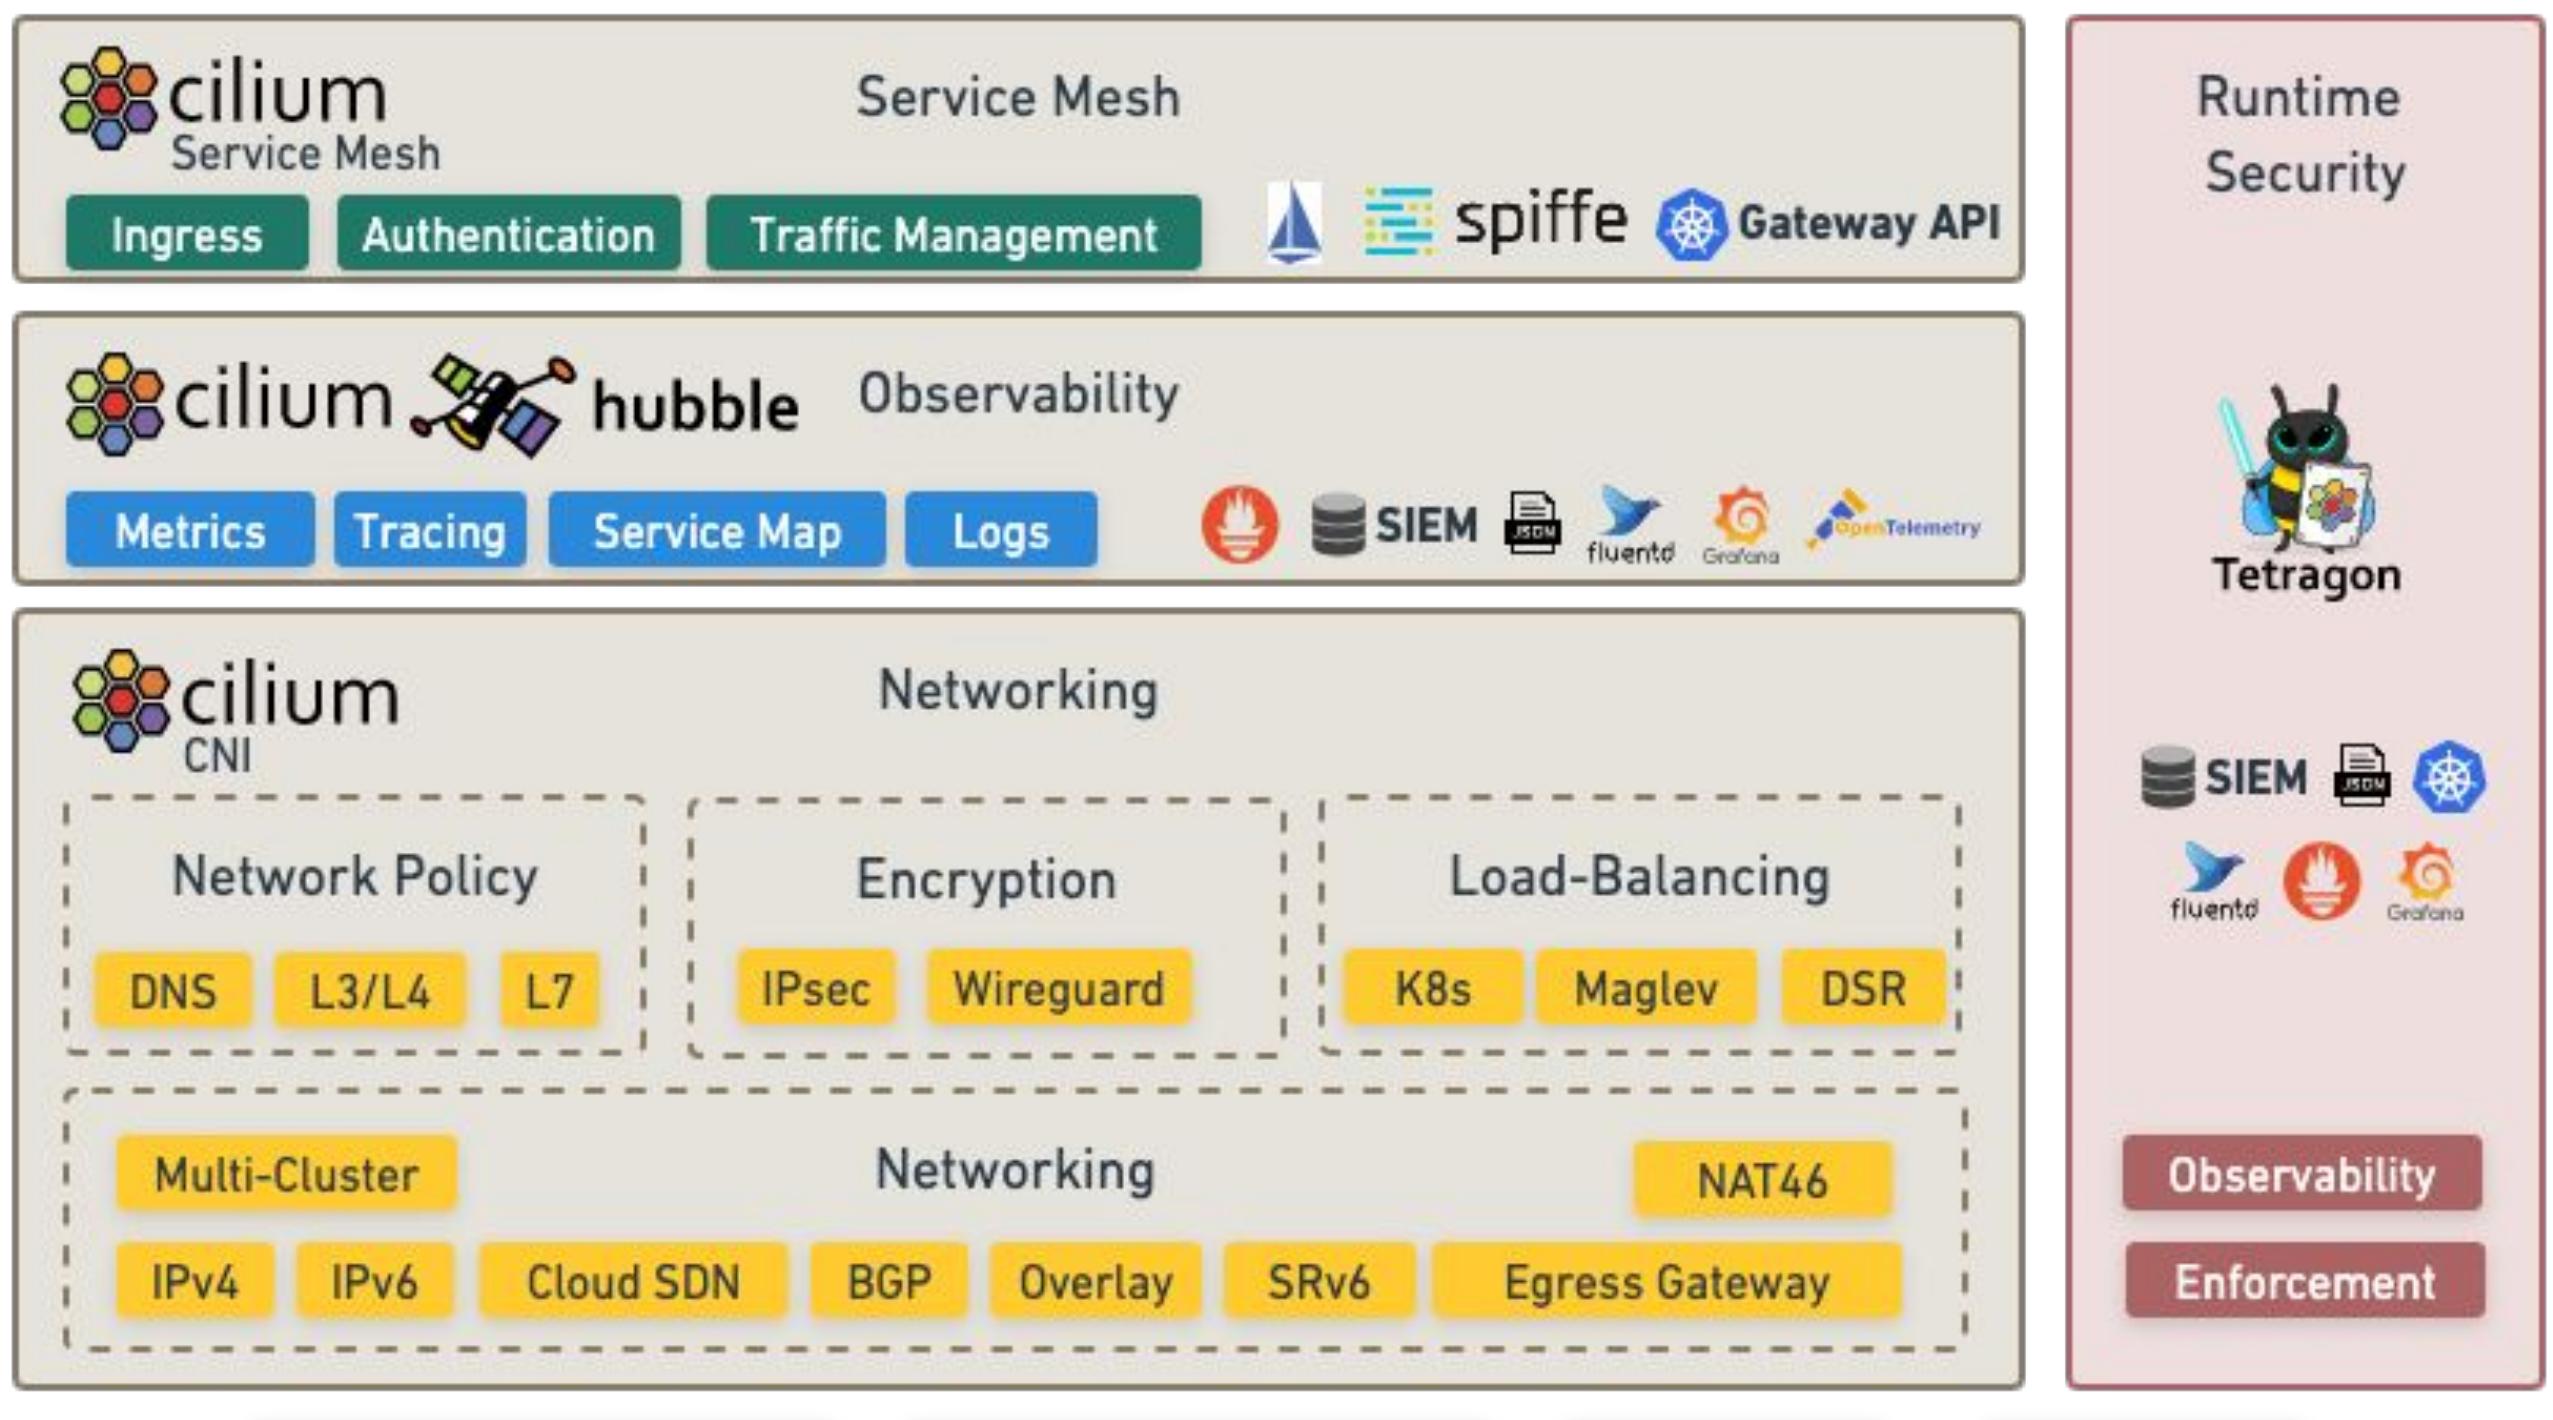
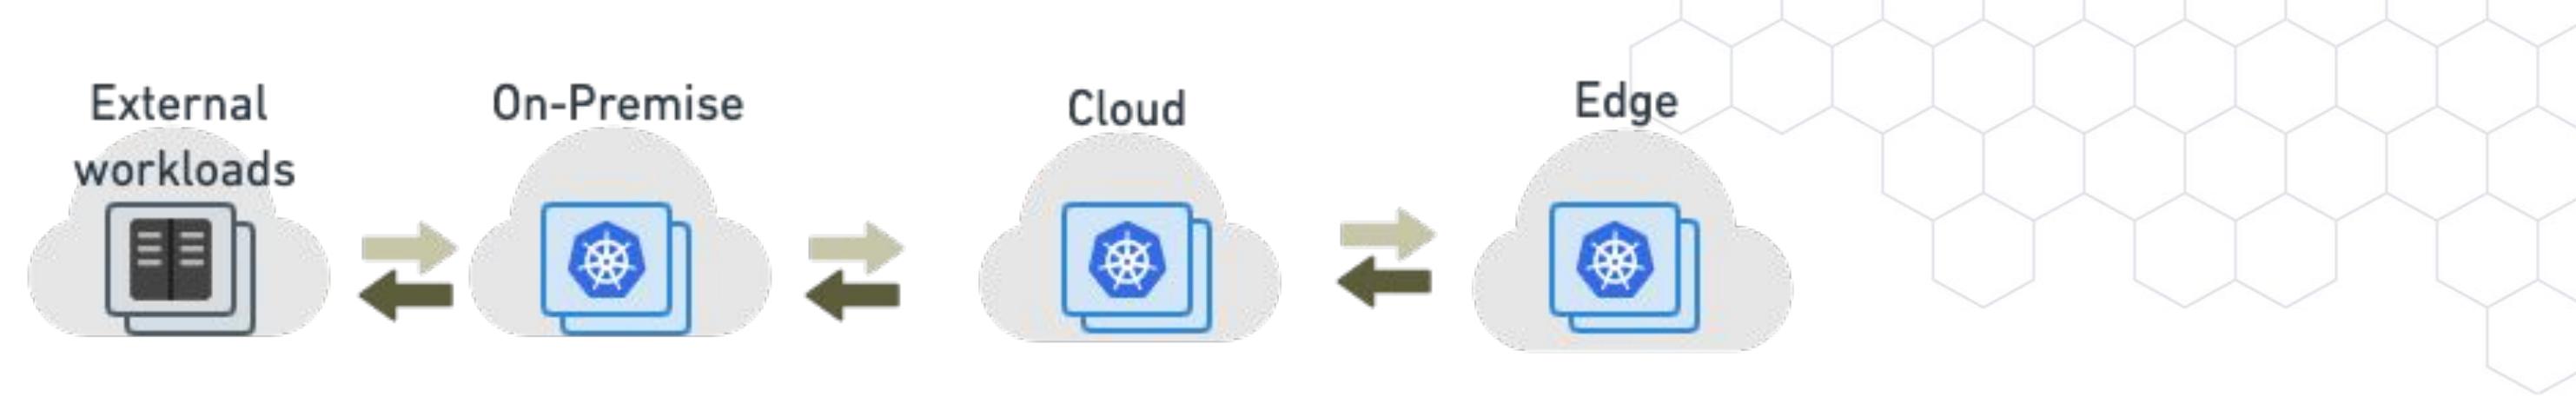
What is Cilium?

- **Networking & Load-Balancing**
 - CNI, Kubernetes Services, Multi-cluster, VM Gateway
- **Network Security**
 - Network Policy, Identity-based, Encryption
- **Observability**
 - Metrics, Flow Visibility, Service Dependency

At the foundation of Cilium is the new Linux kernel technology eBPF, which enables the dynamic insertion of powerful security, visibility, and networking control logic within Linux itself. Besides providing traditional network level security, the flexibility of BPF enables security on API and process level to secure communication within a container or pod.

[Read More](#)







cilium

Created by ISOVALENT

 eBPF-based:

- Networking
- Security
- Observability
- Service Mesh & Ingress

Foundation

 CLOUD NATIVE COMPUTING FOUNDATION

Technology



Building a Global Multi Cluster Gaming Infrastructure with Cilium



What Makes a Good Multi-tenant Kubernetes Solution



Building a Secure and Maintainable PaaS



Building High-Performance Cloud-Native Pod Networks



Cloud Native Networking with eBPF



Managed Kubernetes: 1.5 Years of Cilium Usage at DigitalOcean



Scaling a Multi-Tenant k8s Cluster in a Telco



First step towards cloud native networking



Google chooses Cilium for Google Kubernetes Engine (GKE) networking



Why eBPF is changing the Telco networking space?



Kubernetes Network Policies in Action with Cilium



AWS picks Cilium for Networking & Security on EKS Anywhere



Scaleway uses Cilium as the default CNI for Kubernetes Kapsule



Sportradar is using Cilium as their main CNI plugin in AWS (using kops)



Utmost is using Cilium in all tiers of its Kubernetes ecosystem to implement zero trust



Yahoo is using Cilium for L4 North-South Load Balancing for Kubernetes Services

ISOVALENT



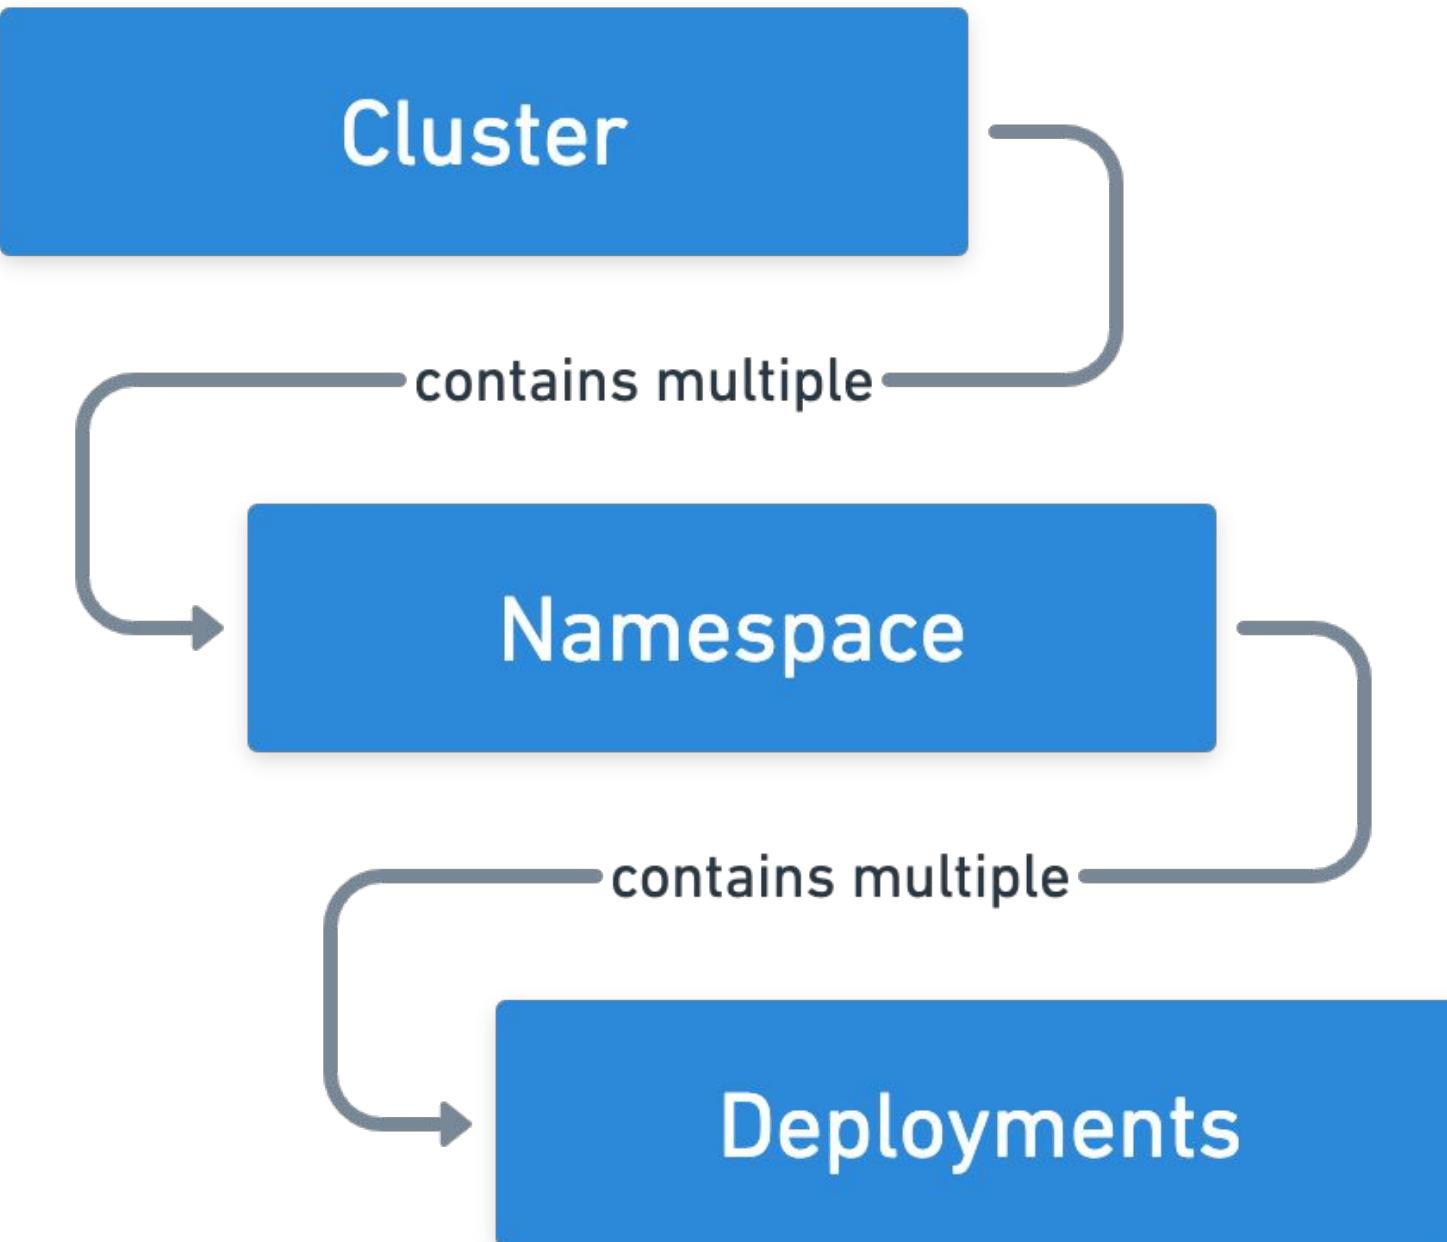
Zero Trust Networking

Using Cilium, Hubble & Grafana

- Cilium & eBPF
- Network Policies Strategies
- Network Policies Observability

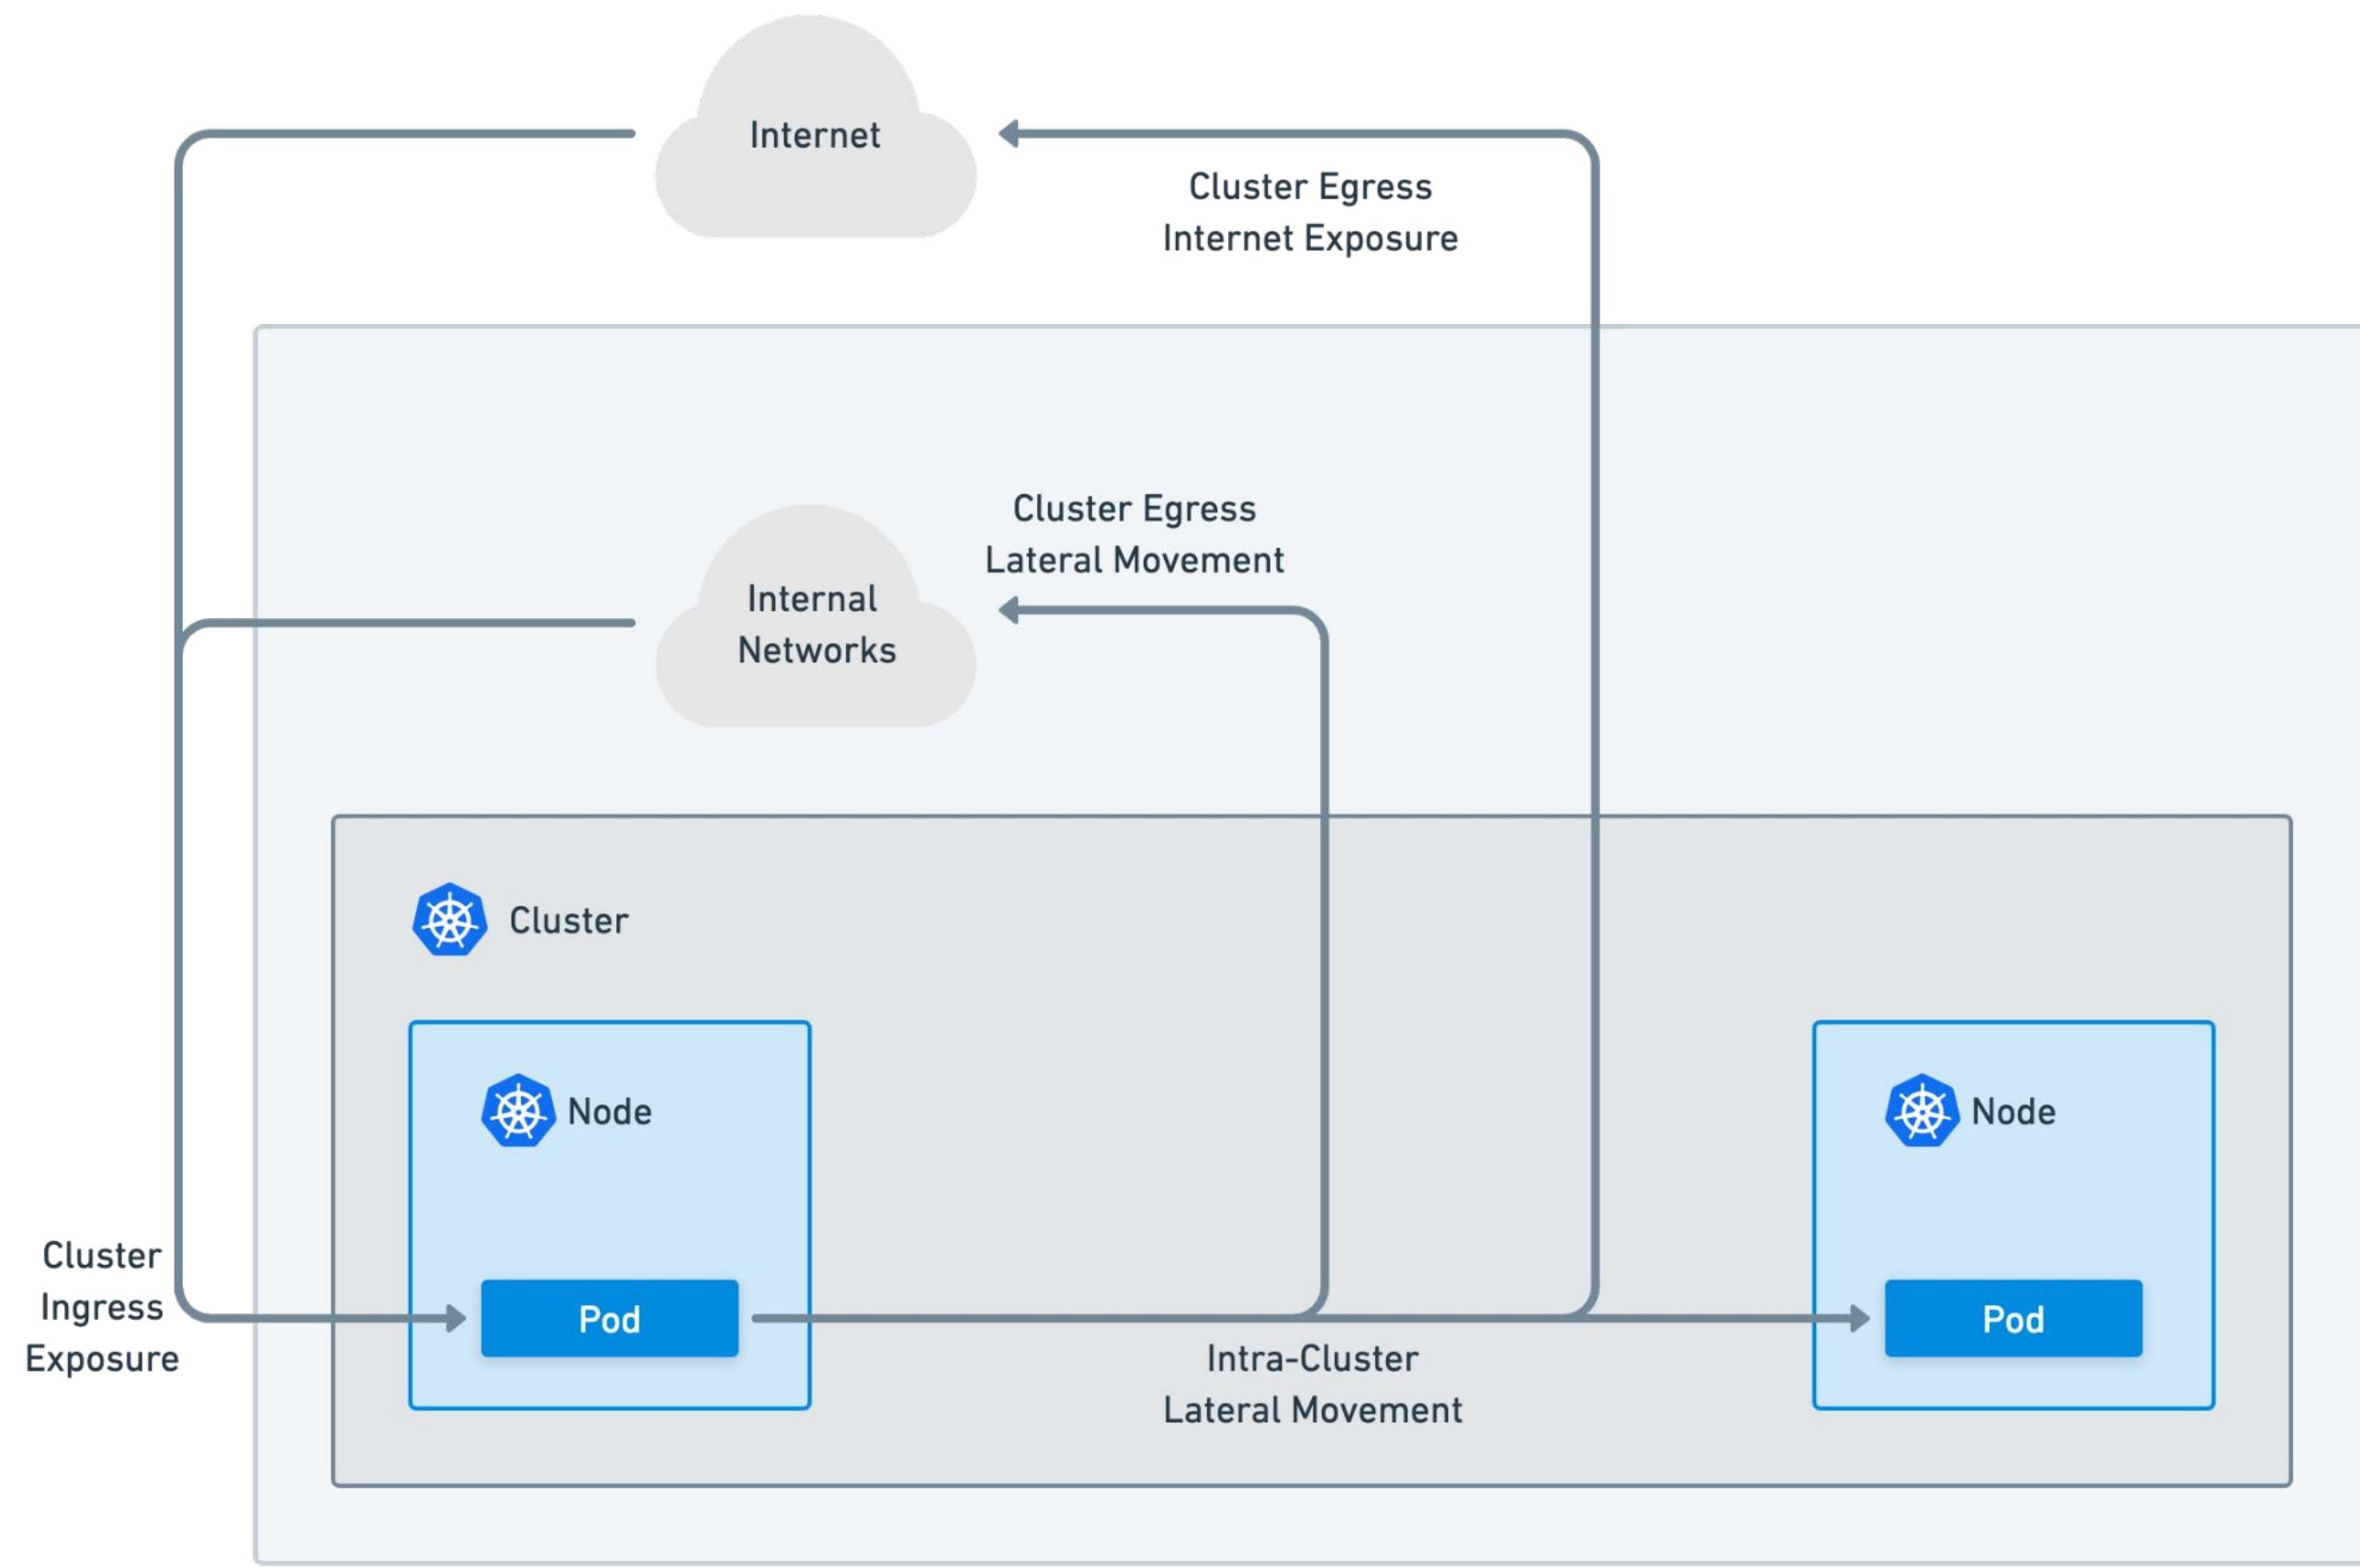
Application Teams & Multi-tenancy

- A Kubernetes cluster is often used by multiple teams, each managing one or more services.
- Each team has one or more namespaces to run their applications.
- Applying network policy at namespace level is a common first step in enabling multi-tenancy.

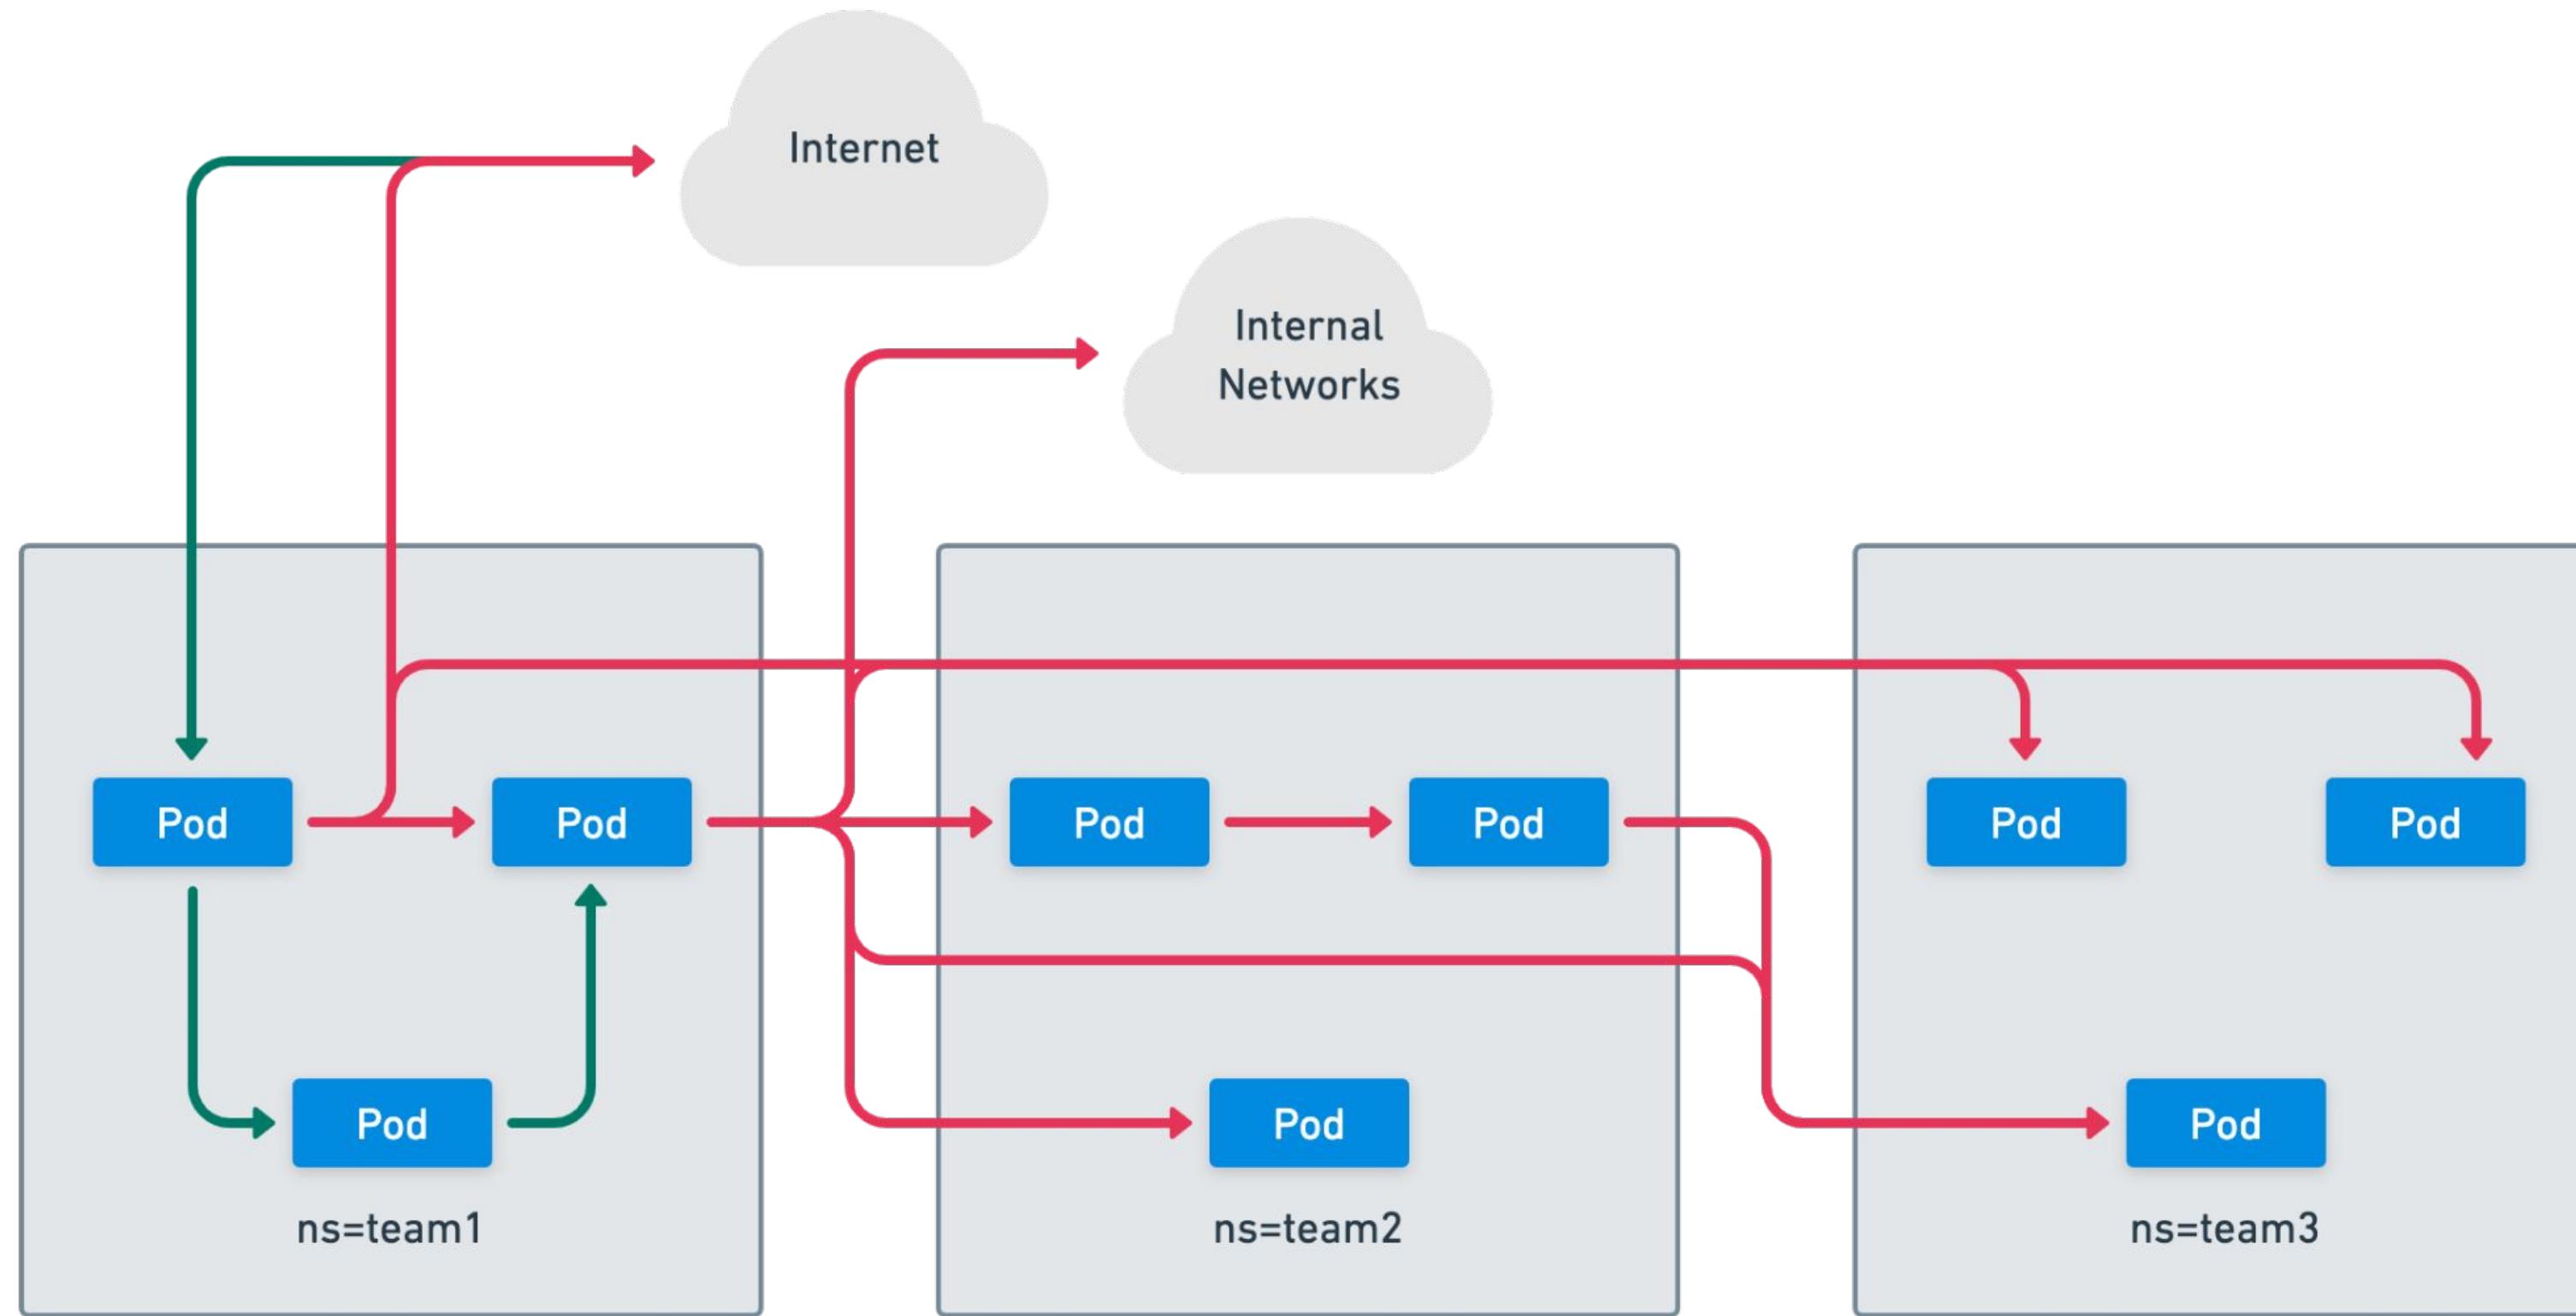


Note: Not always 1-to-1 team-to-namespace. Team may be a namespace-label, or pod label within a larger namespace, but same concept applies.

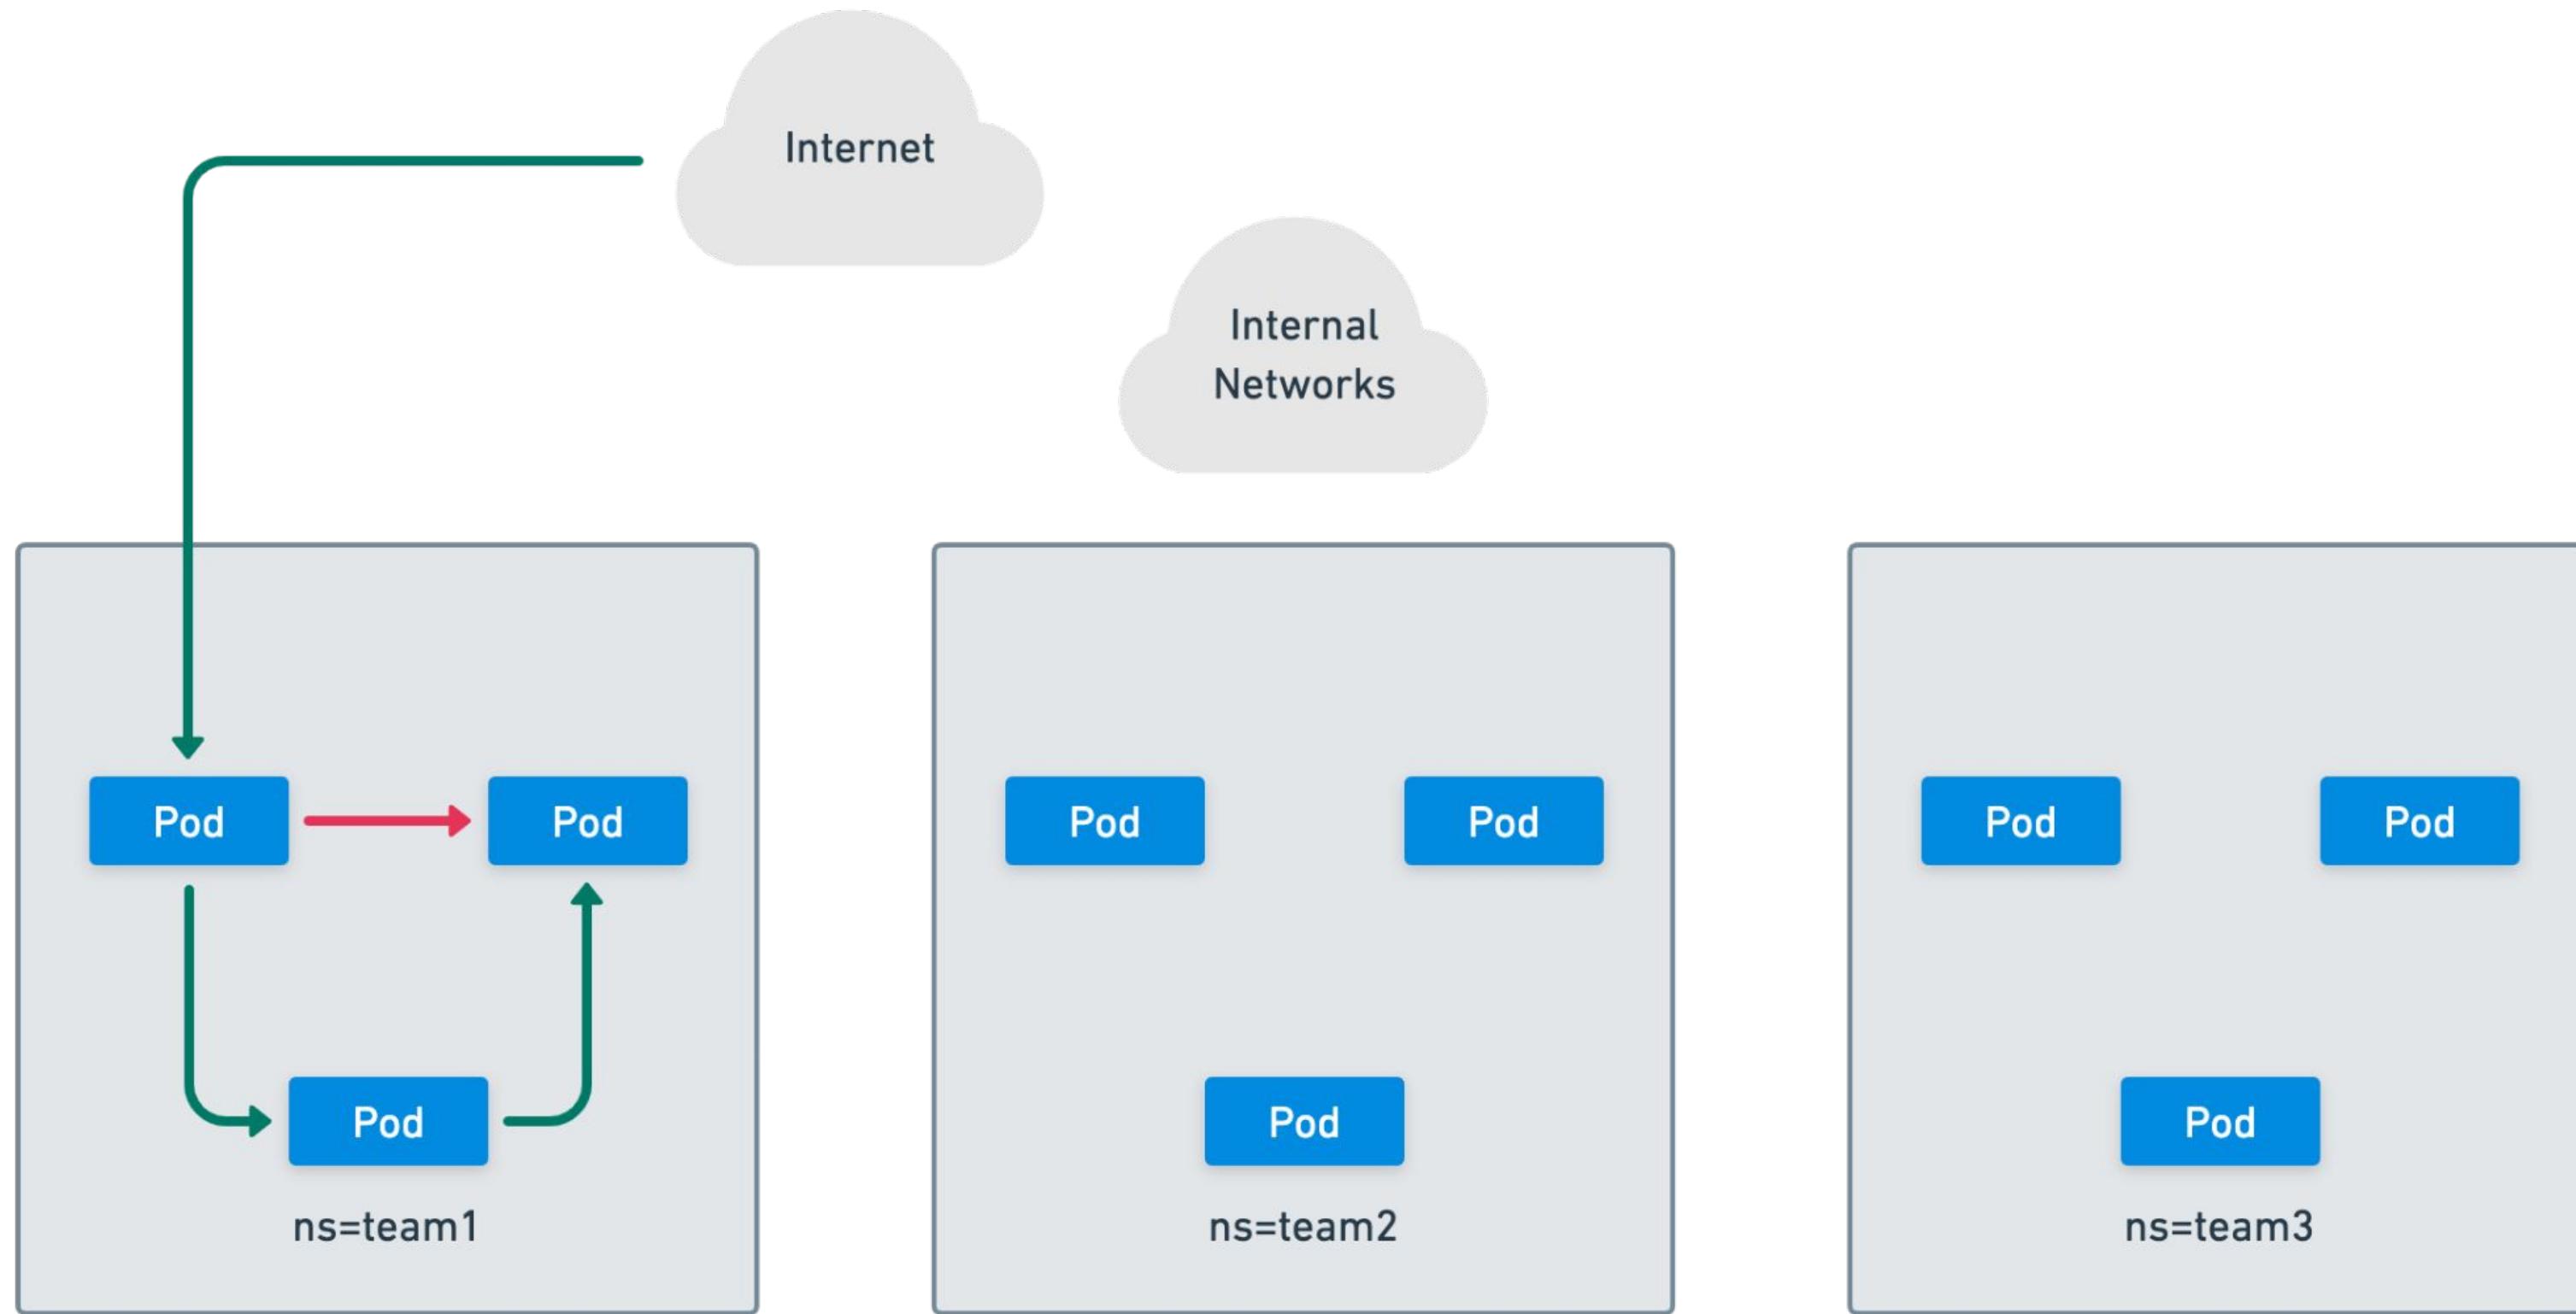
Key Types of Kubernetes Network Risk Exposure



Risk Reduction Example - Wide blast radius



Risk Reduction Example - Limited blast radius





Measuring Risk: Prioritization

Concrete Measurements to guide prioritization

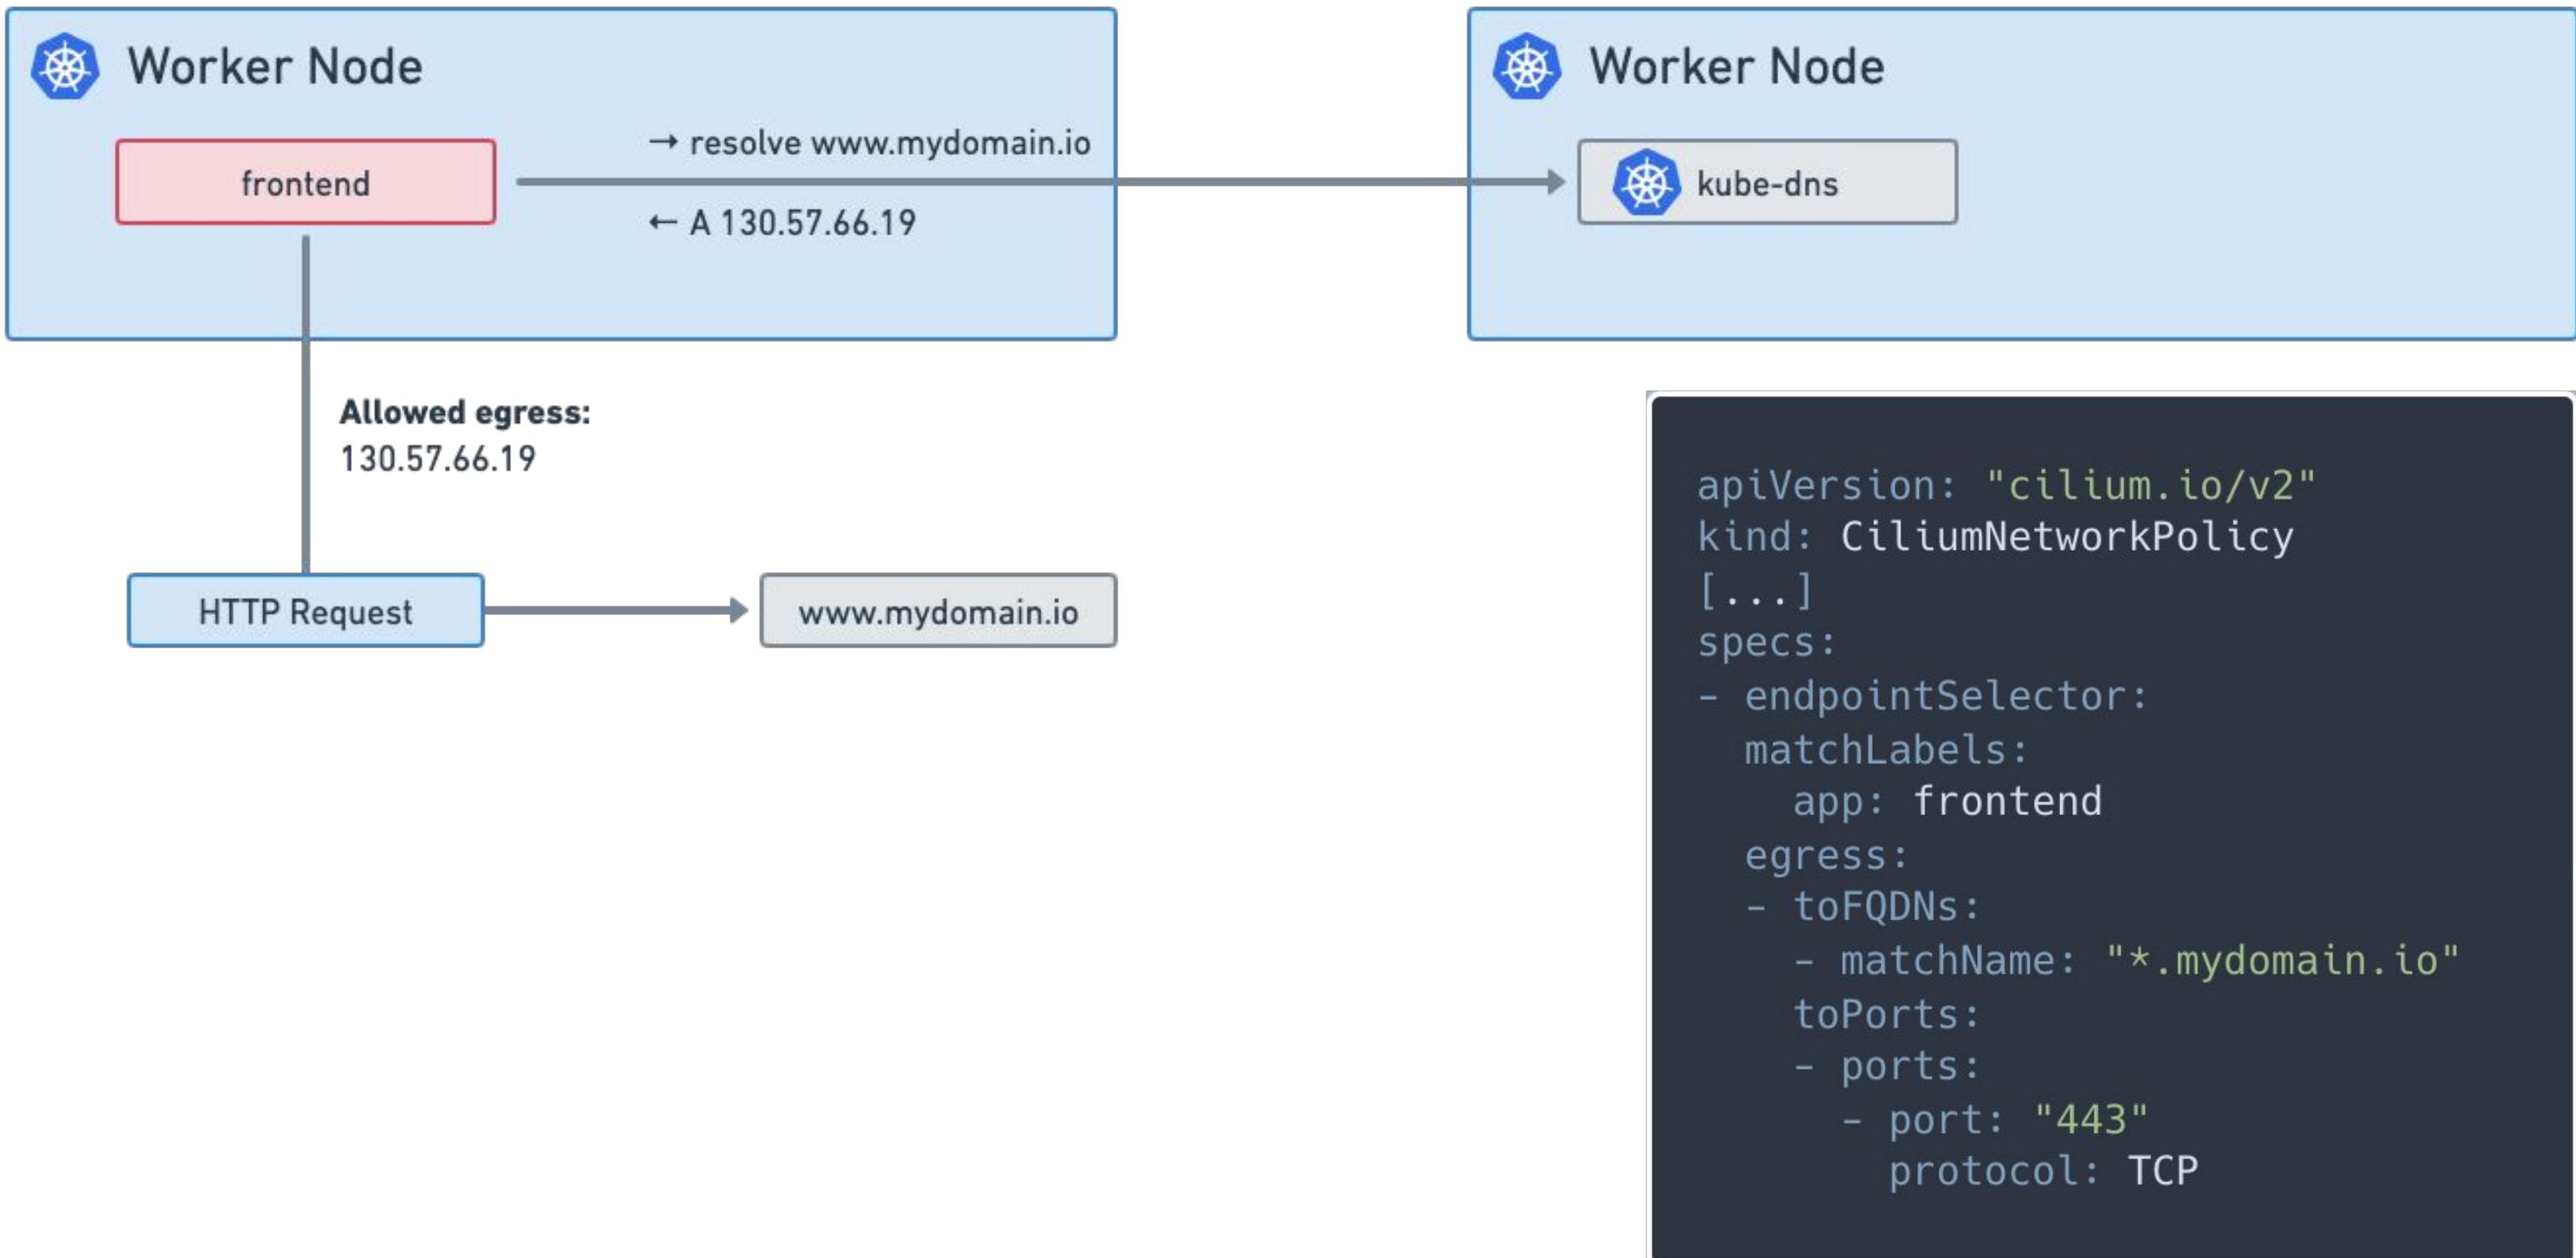
- # of services reachable via Cluster Ingress with no matching traffic.
- # of services reachable from other namespaces with no matching traffic.
- # of services with access to External Networks or Internet with no matching connections.

The Cost of “Overfitting”

- Strong network observability tools enables you to translate each observed flow into a rule.
- However, simplified approaches to translating traffic flow to rule tend to be cumbersome and brittle.
- Why?:
 - IPs for services outside the cluster can change or are opaque.
 - App dependencies or shared services: kube-dns, external logging, vaults servers are required by all apps. Creating pair-wise rules leads to friction and bloat.
 - Communication within a single app/team. Complex and frequent change. Leads to increased number of rules which change frequently.

Best Practices - DNS Rules

DNS-aware Cilium Network Policies for External Services which IPs frequently change.





Initial “Coarse-Grained” Per-Namespace Strategy

Pattern to avoid “overfitting”

1. **Allow all** ingress/egress communication within a namespace

2. Use **Hubble Observability** data to identify and permit:
 - a. Which (service + port) within the namespace are **“public services”**:
 - Allow from cluster-only.
 - Allow from “world” (type LB/NP) or ingress namespace.
 - b. **Egress** Access to:
 - Other services in the cluster (optional, limit ports).
 - Other services in the external private network (optional, limit ports).
 - Other services on the Internet (optional, limit ports).

3. **Coarse-grained** nature of policies ⇒ policies need only change **rarely**

Key Pattern: Baseline vs. Per-Namespace Policies



Global Baseline Policies

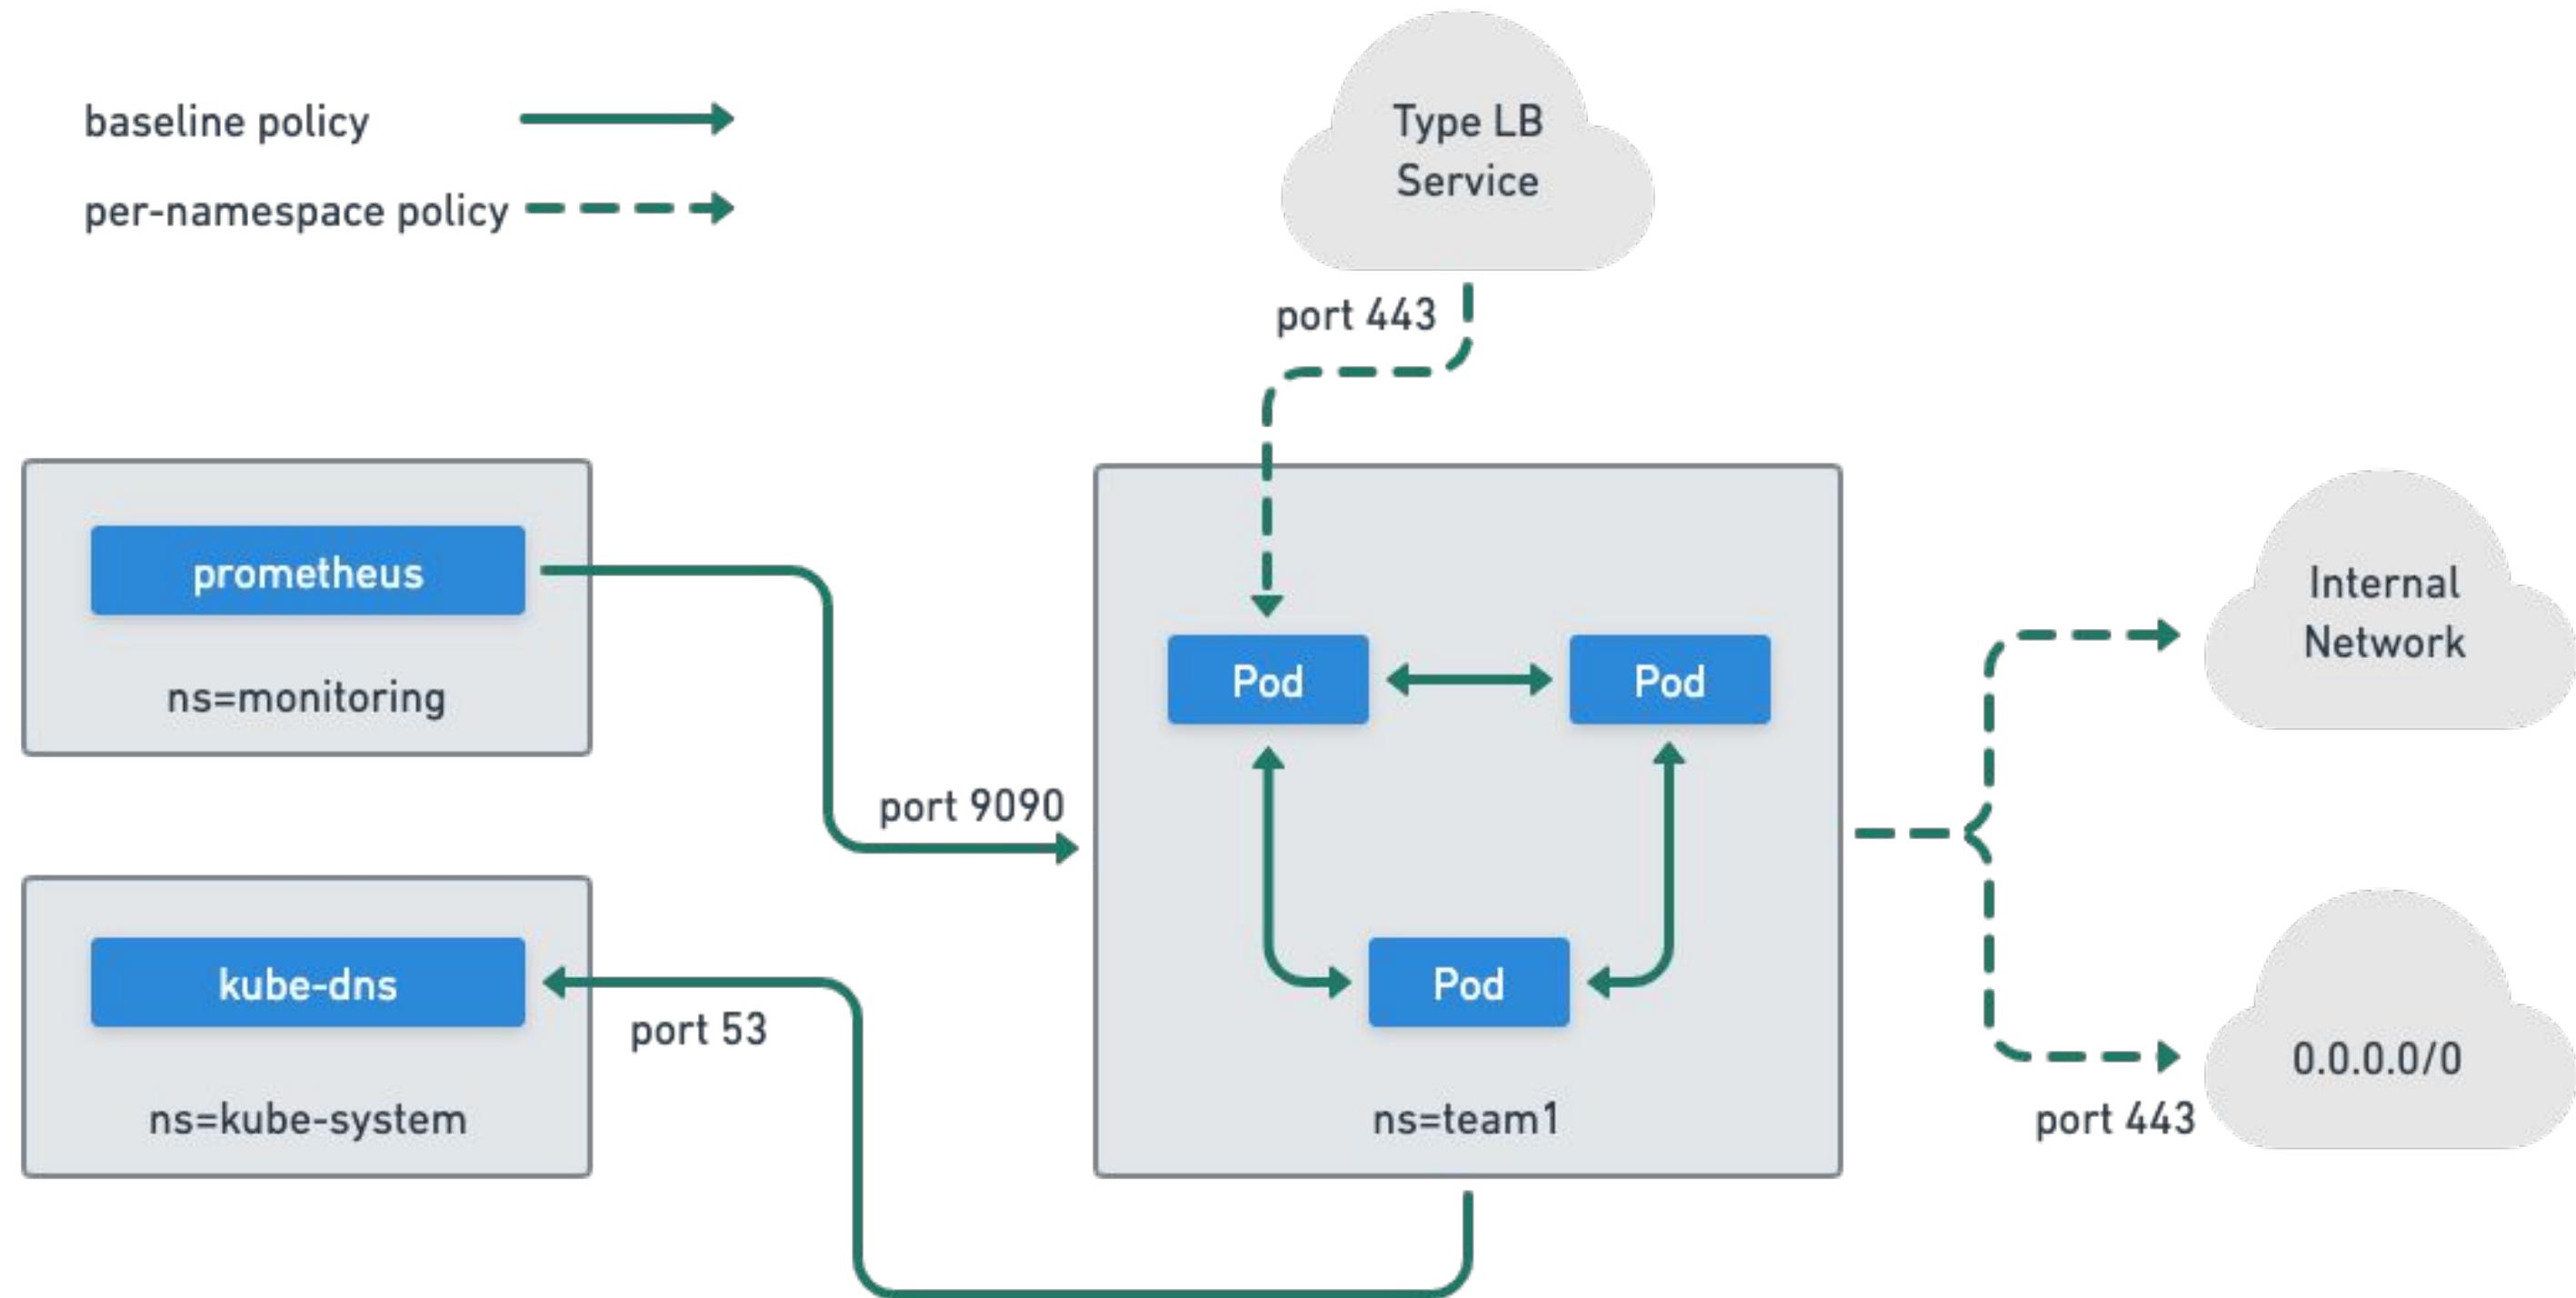
- Default deny ingress/egress.
- Allow all ingress/egress within namespace.
- Egress to “public services”
 - cluster-wide shared services (e.g. kube-dns, prometheus).
 - external shared CIDR/DNS (logging, monitoring, vault, etc.)
- Often implemented by

CiliumClusterWideNetworkPolicies

Per-Namespace Policies

- Per-service ingress (limited by port):
 - Exposed within cluster via ClusterIP service.
 - Exposed externally via ingress or Load-Balancer/NodePort.
- Namespace-specific egress:
 - No access.
 - Egress to specific CIDR/DNS + port
 - Unrestricted access on specific ports (fallback).

Example Baseline + Coarse-grained Per Namespace Policies

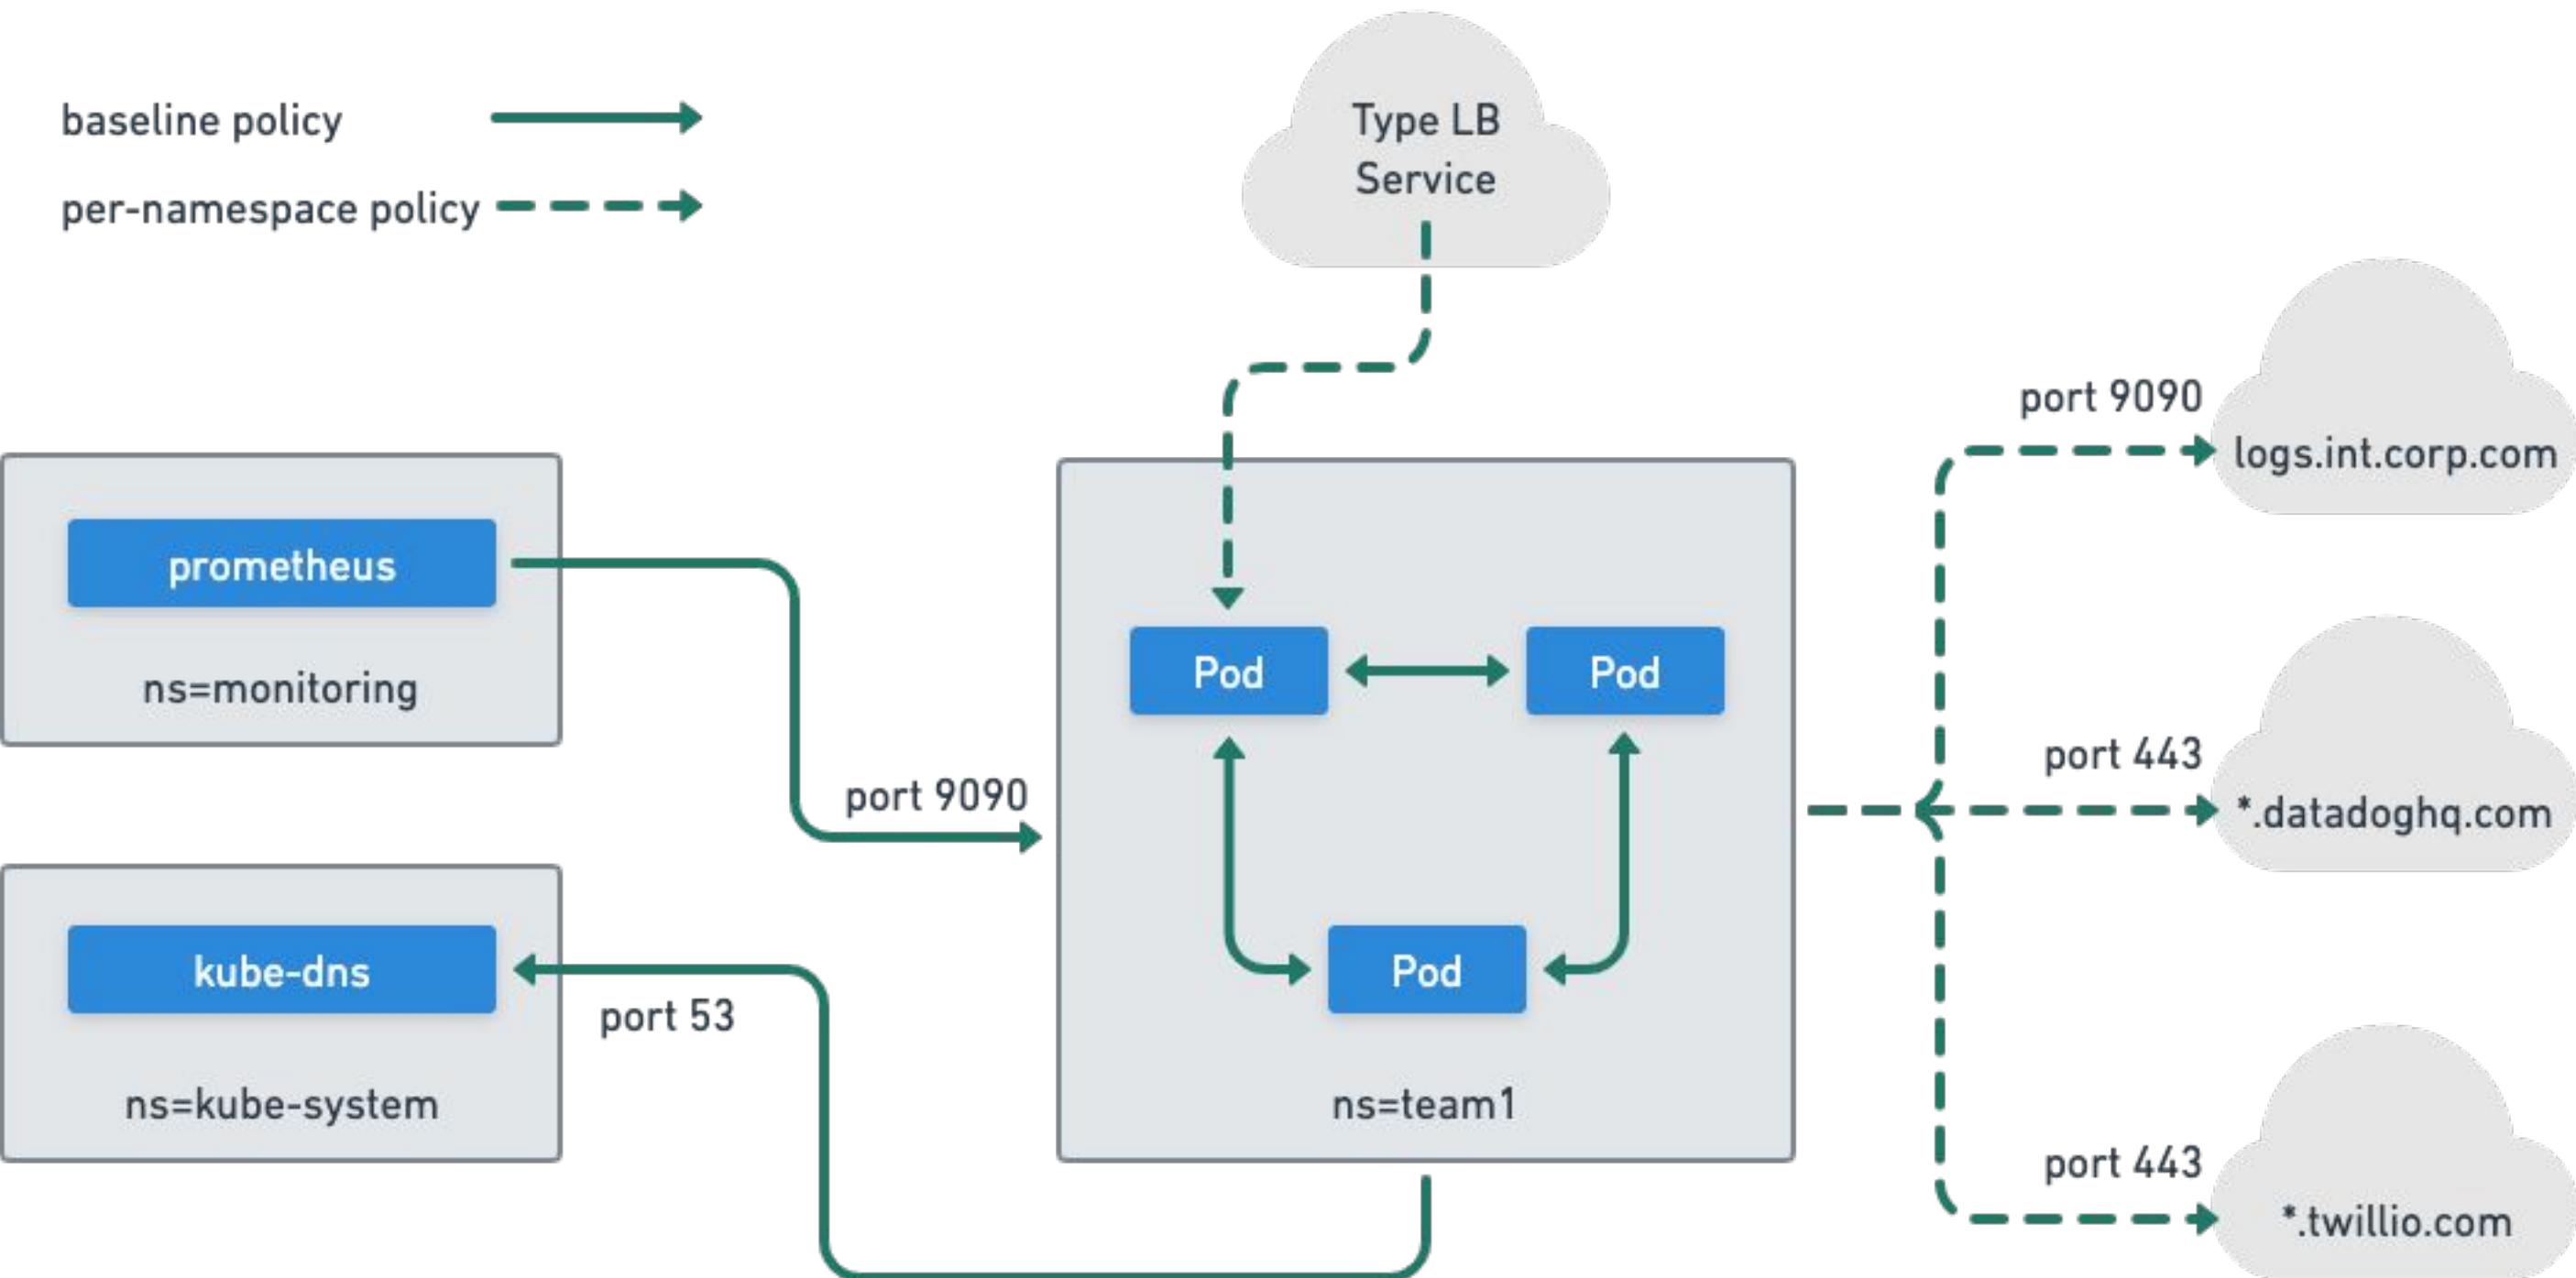


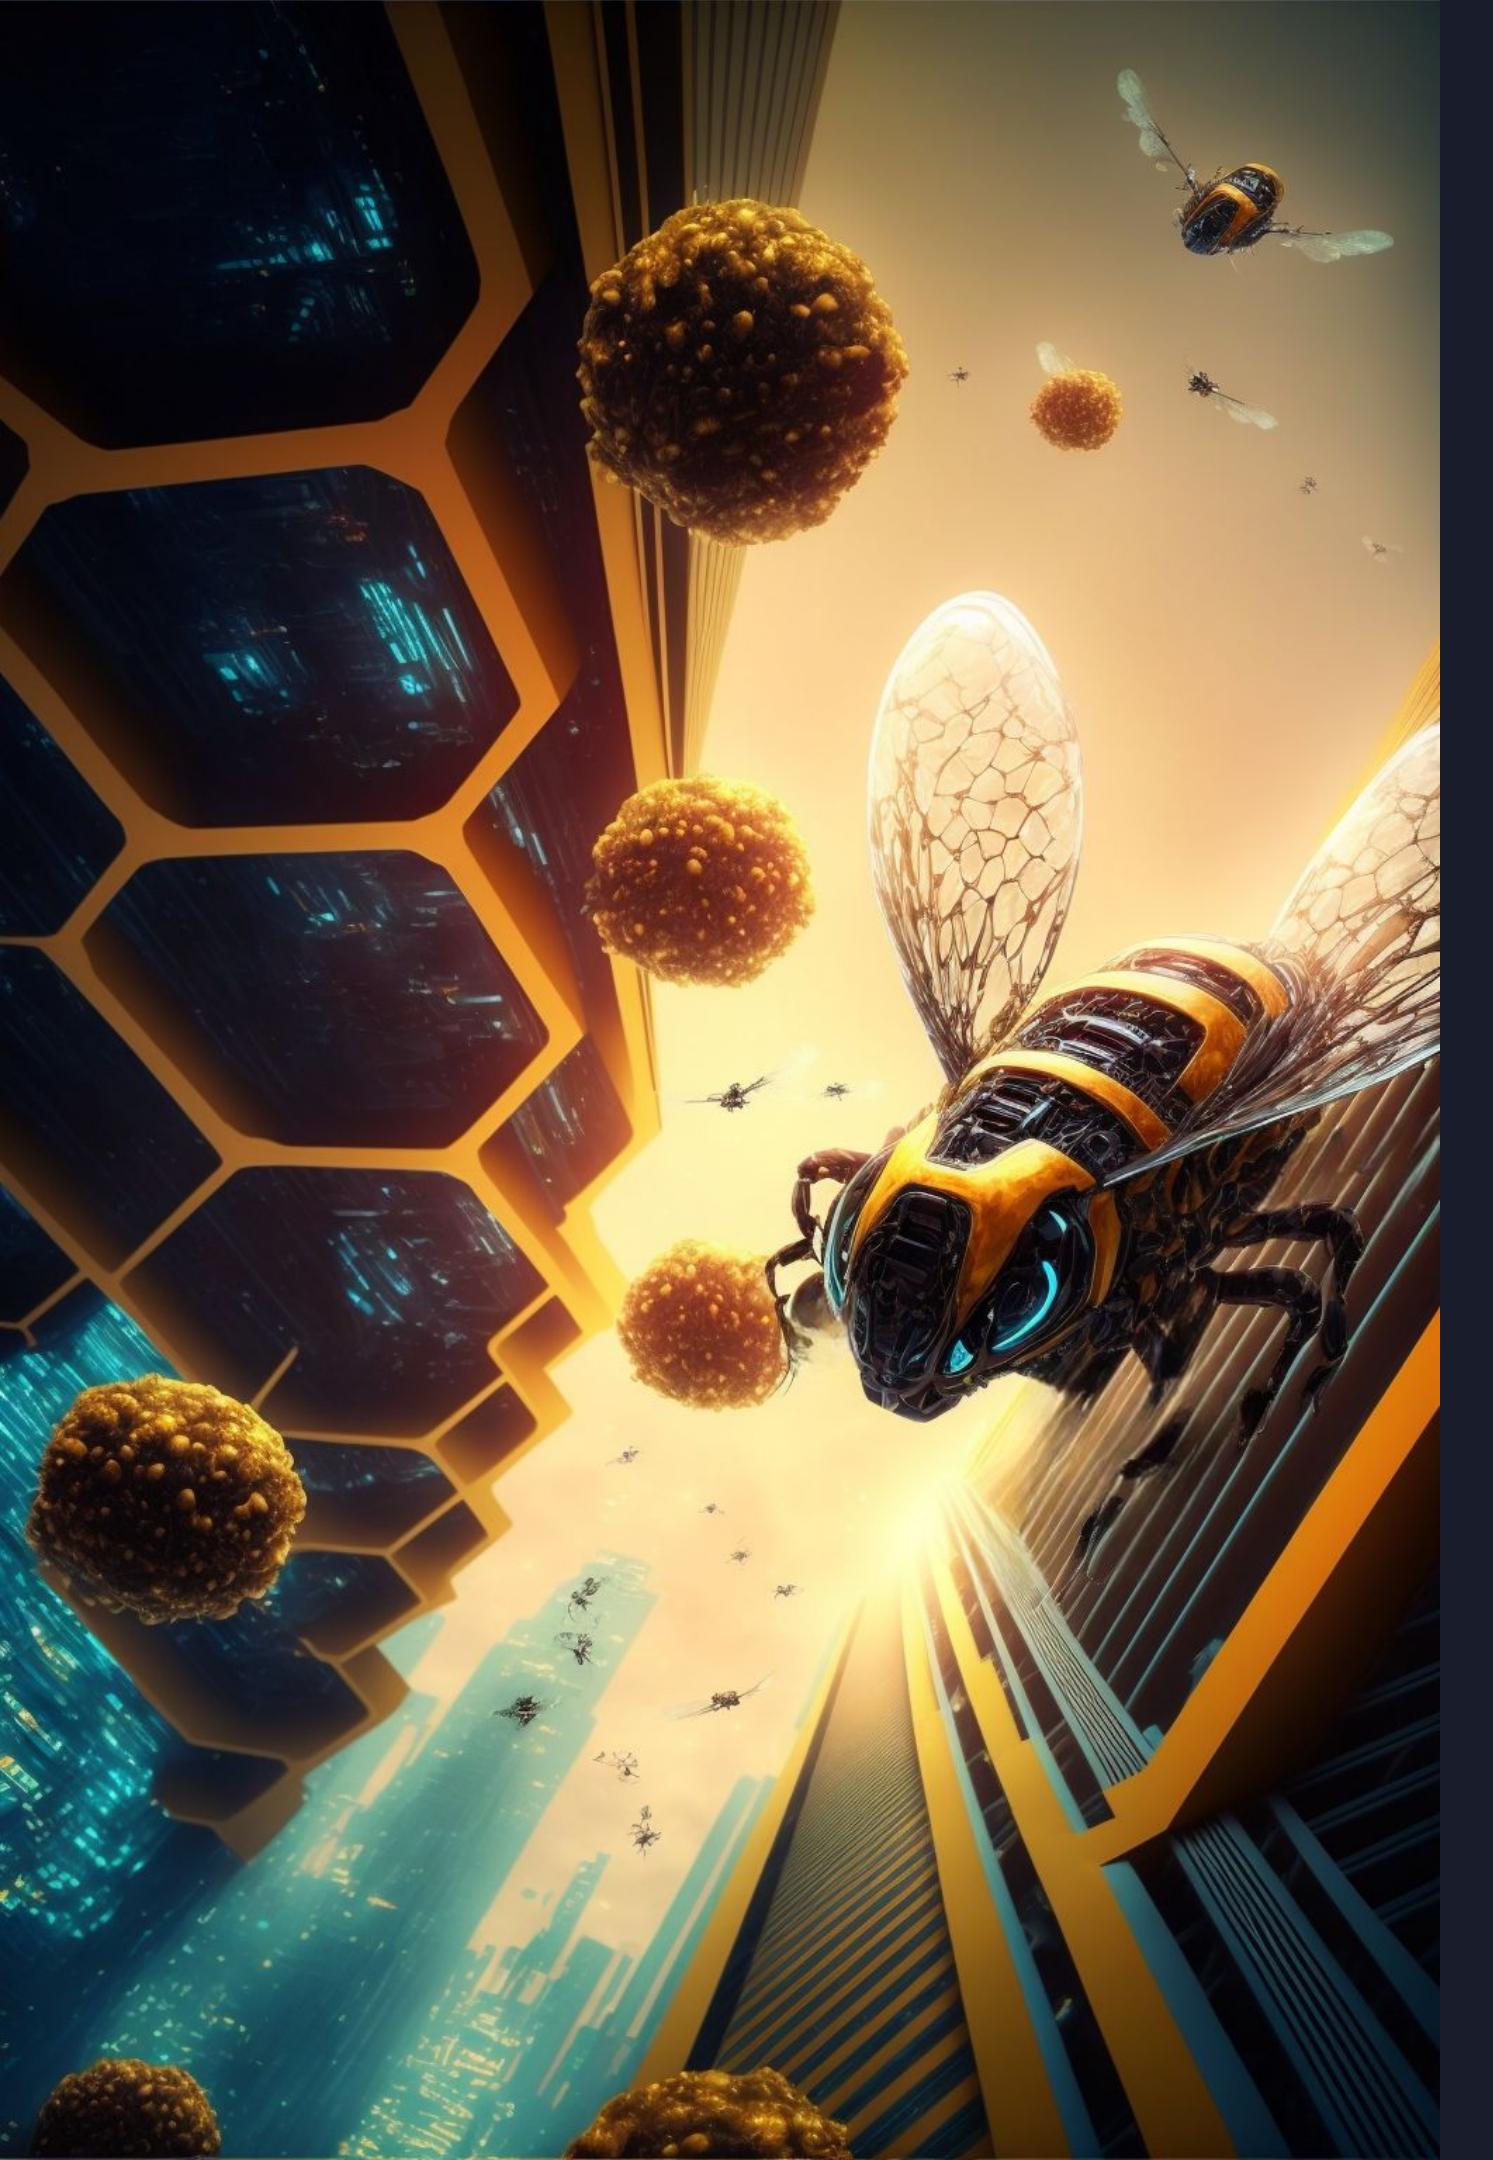


Transitioning from “Coarse” to “Fine-Grained” Policies

- **Prioritize namespaces based on:**
 - Most **security-sensitive** applications
 - Measurement of exposure vs. used connectivity
- **Transition for egress** outside the cluster:
 - Access to all private network → Access to specific FQDNs or smaller CIDRs (with port)
 - Access to Internet → Access to specific FQDNs or small CIDRs (with port)
- **Transitions within** the cluster:
 - Shift rules that allow all to/from cluster to allowing to/from specific namespaces or services.

Example Baseline + Fine-grained Per Namespace Policies



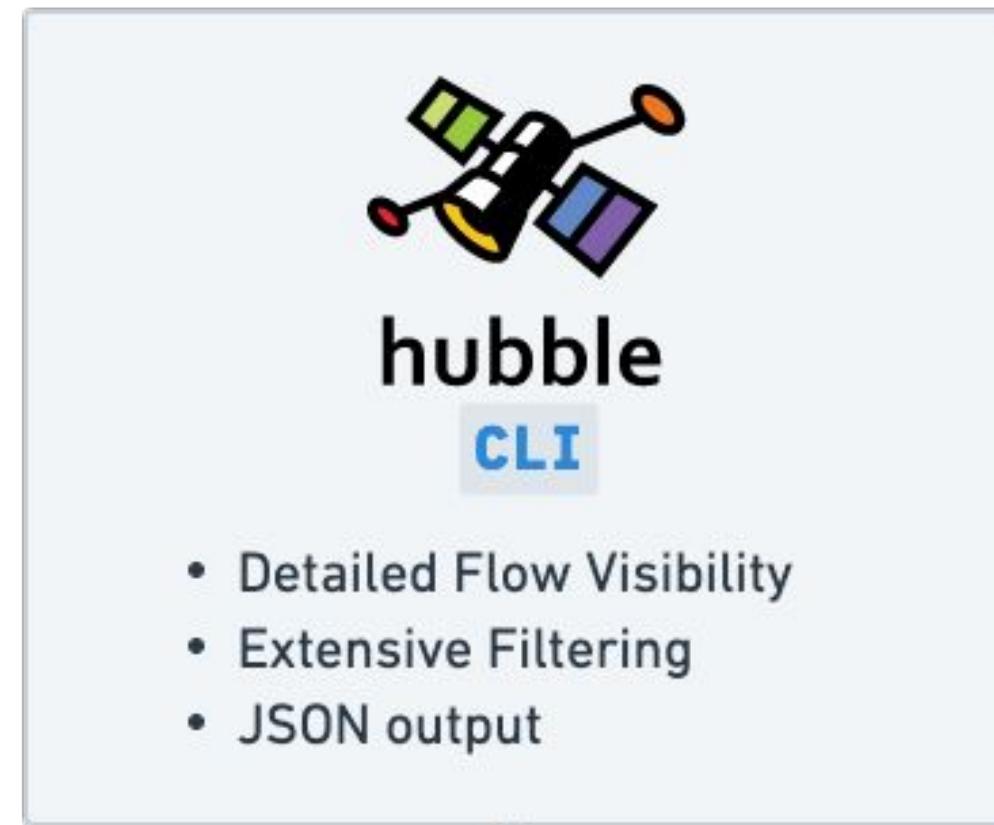
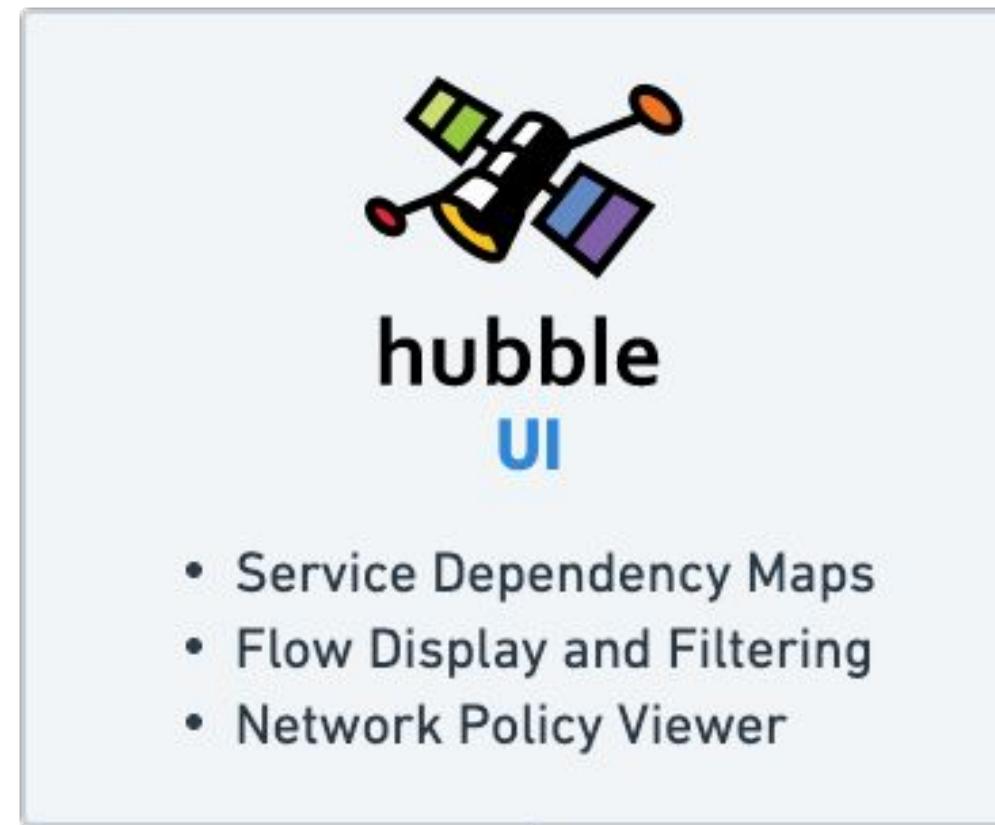


Zero Trust Networking

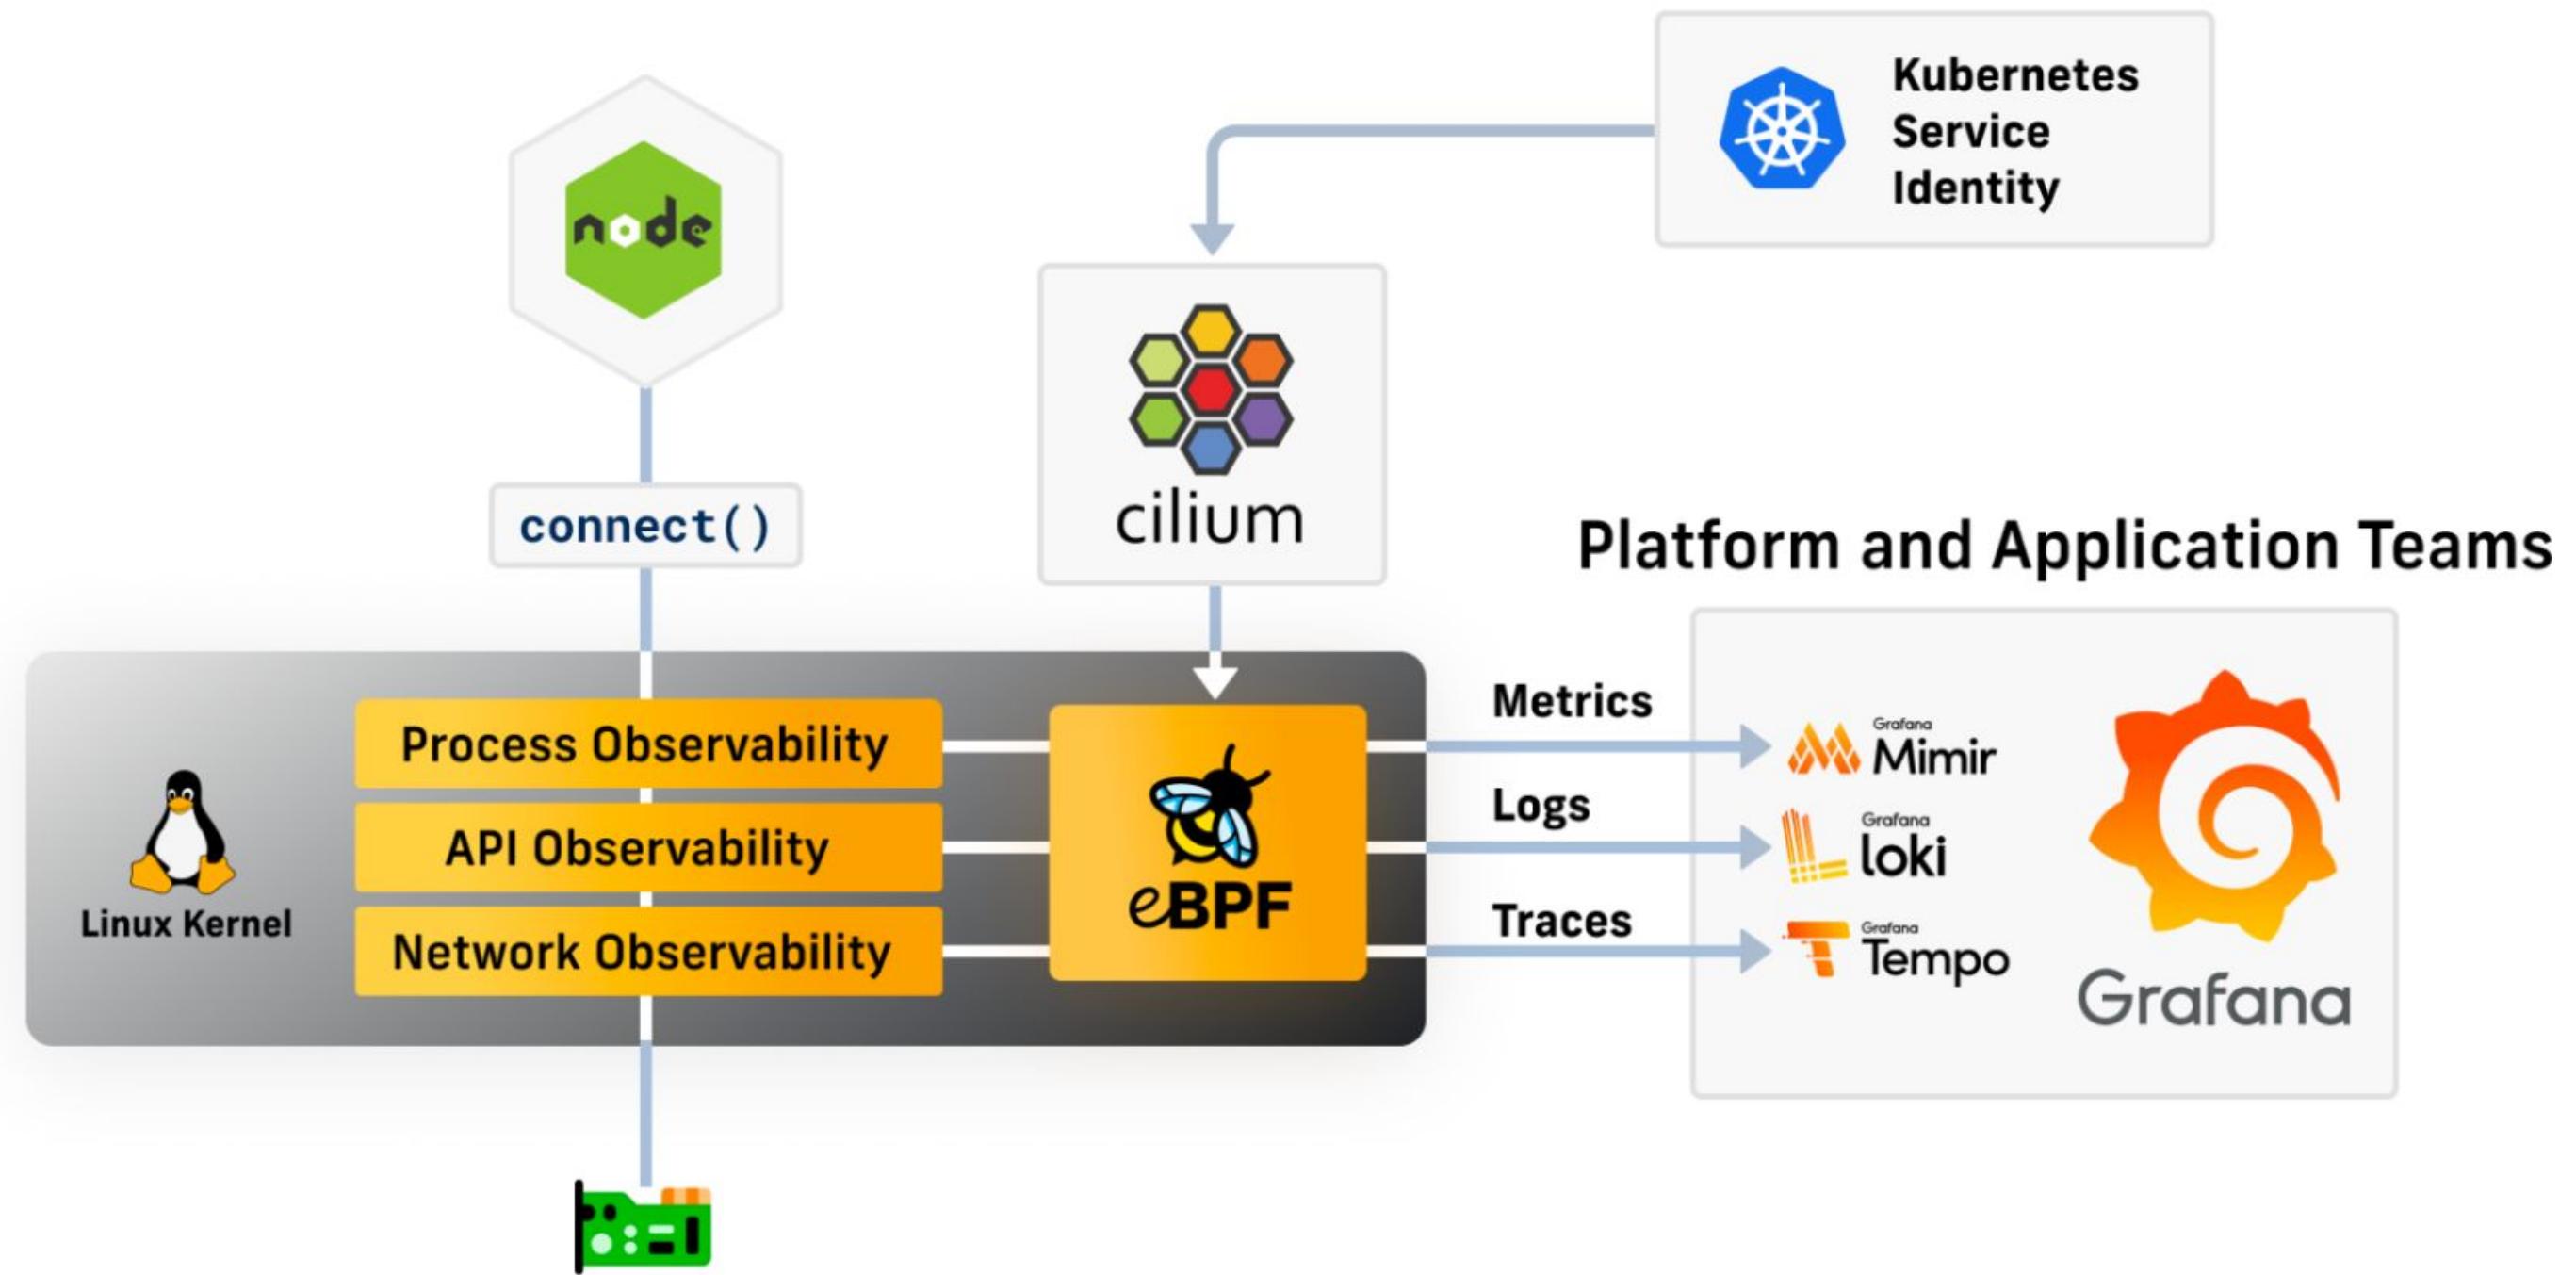
Using Cilium, Hubble & Grafana

- Cilium & eBPF
- Network Policies Strategies
- Network Policies Observability

Hubble Overview



Cilium + Hubble ❤️ Grafana



Using Hubble Observability Data to Build Policies

Filter by: label key=val, ip=1.1.1.1, dns=google.com, identity=42, pod=frontend

The Service Map displays the network topology of the 'jobs-app' namespace. Services shown include crawler, recruiter, loader, jobposting, coreapi, elasticsearch, kafka, and zookeeper. External services like api.twitter.com and Unknown App are also connected. The map shows bidirectional traffic between these components.

Live View

Namespace: jobs-app

Flows verdict:

- Any verdict
- Forwarded
- Dropped

Aggregate flows

Visual filters:

- Host service
- Kube-DNS:53 pod
- Remote node
- Prometheus app

Notifications: 2.4K flows/s • 5/5 nodes

raymond.dejong@isovalent.com

Source Identity	Destination Identity	Destination Port	L7 info	Verdict	TCP Flags	Timestamp
crawler jobs-app	api.twitter.com	443	—	forwarded	SYN	2022/10/05 11:28:27 (+02)
crawler jobs-app	api.twitter.com	443	—	forwarded	SYN	2022/10/05 11:28:22 (+02)
crawler jobs-app	api.twitter.com	443	—	forwarded	SYN	2022/10/05 11:28:17 (+02)
crawler jobs-app	api.twitter.com	443	—	forwarded	SYN	2022/10/05 11:28:12 (+02)
crawler jobs-app	loader jobs-app	50051	—	forwarded	ACK PSH	2022/10/05 11:28:12 (+02)
jobposting jobs-app	coreapi jobs-app	9080	—	forwarded	SYN	2022/10/05 11:27:58 (+02)
No app name	kafka jobs-app	9092	—	forwarded	ACK PSH	2022/10/05 11:27:57 (+02)
kafka jobs-app	zookeeper jobs-app	2181	—	forwarded	ACK PSH	2022/10/05 11:27:55 (+02)
No app name jobs-app	kafka jobs-app	9092	→ Kafka	forwarded		2022/10/05 11:27:52 (+02)
coreapi jobs-app	elasticsearch jobs-app	9200	→ GET /jobs/_search 0ms	forwarded		2022/10/05 11:27:38 (+02)
jobposting jobs-app	coreapi jobs-app	9080	→ GET /jobs 0ms	forwarded		2022/10/05 11:27:38 (+02)
jobposting jobs-app	coreapi jobs-app	9080	—	forwarded	SYN	2022/10/05 11:27:38 (+02)

Hubble Network Policy Editor

Filter by: label key=val, ip=1.1.1.1, dns=google.com, identity=42, pod=frontend

Dashboard

Service Map

Network Policies

Process Tree

Live View

Namespace

jobs-app

Flows verdict

Any verdict

Forwarded

Dropped

Aggregate flows

Network policies

Visualize all

Selected policy visualized on map

allow-all-within-namespace

fqdn

l7-ingress-visibility

twitter-fqdn

Notifications

2.4K flows/s • 5/5 nodes

raymond.dejong@isovalent.com

Outside Cluster +
Any endpoint

In Namespace jobs-app +
Any pod

In Cluster +
Everything in the cluster

In Namespace jobs-app +
k8s:app=crawler

Outside Cluster +
Any endpoint
api.twitter.com :443|TCP

In Namespace jobs-app +
Any pod

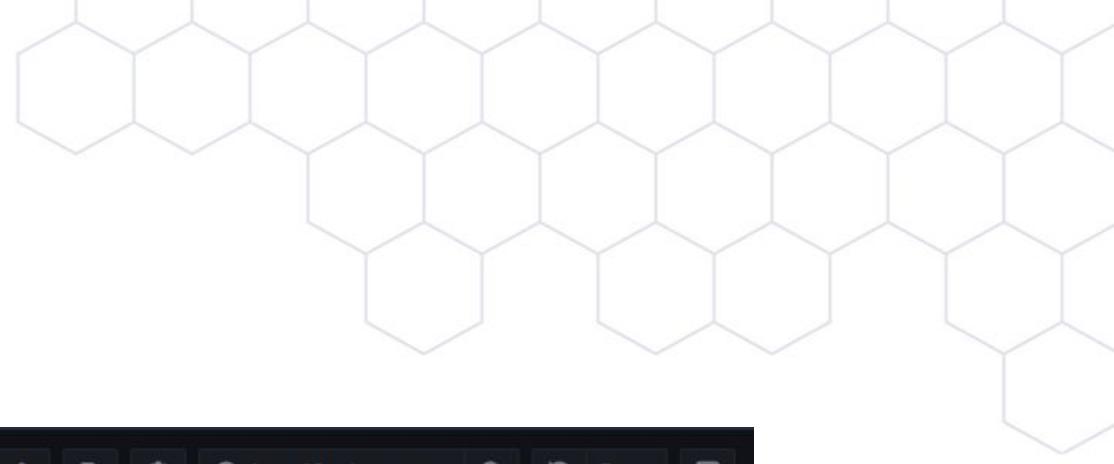
In Cluster +
Everything in the cluster
Kubernetes DNS DNS proxy on

Kubernetes Cilium

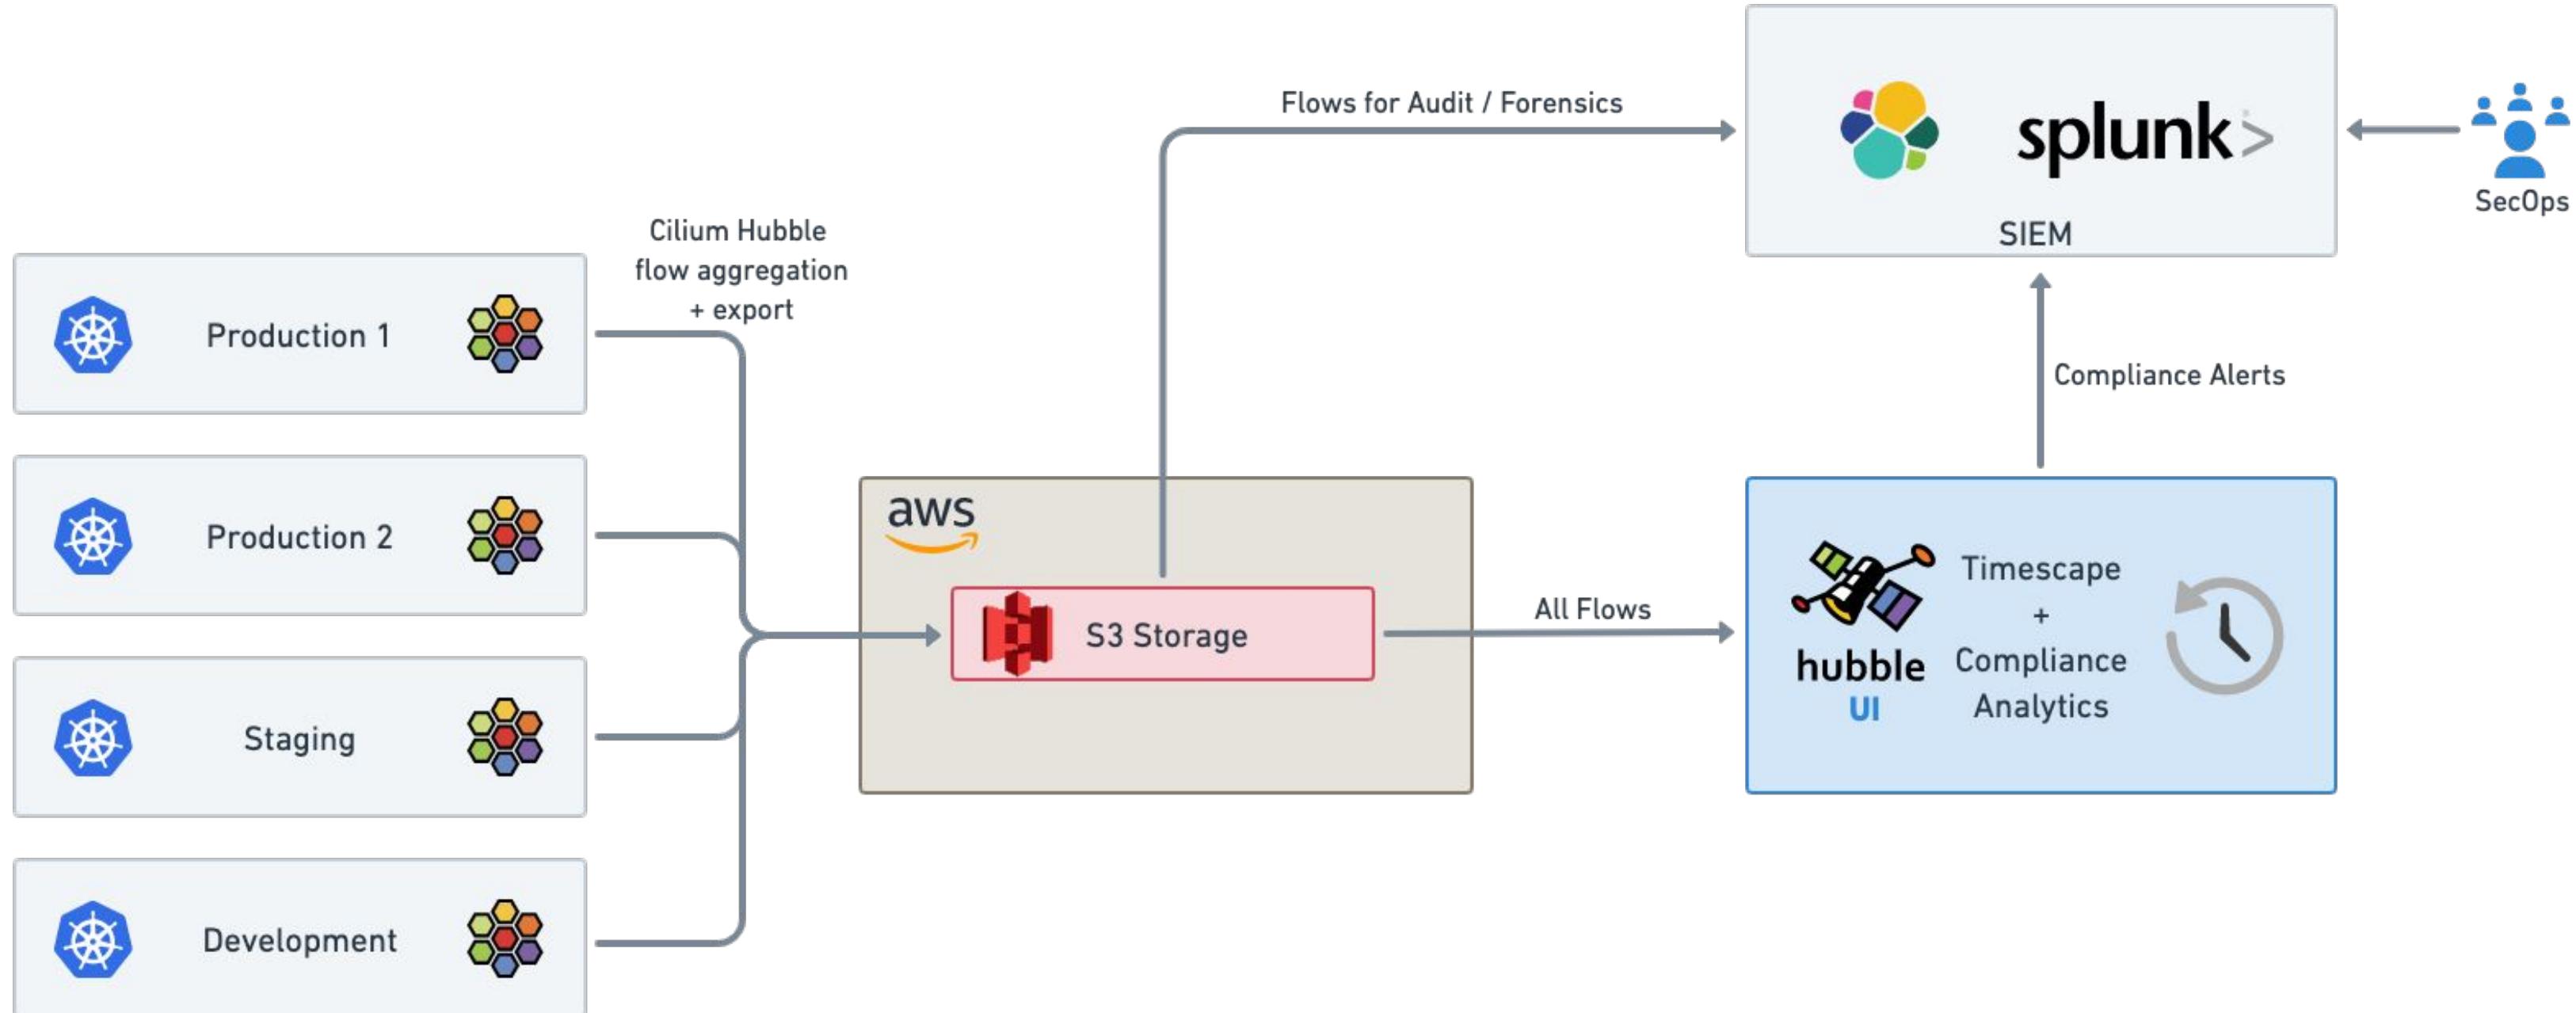
Download

Source Identity	Destination Identity	Verdict
crawler jobs-app	api.twitter.com	forwarded
crawler jobs-app	loader jobs-app	forwarded
jobposting jobs-app	coreapi jobs-app	forwarded
No app name jobs-app	kafka jobs-app	forwarded
No app name	kafka jobs-app	forwarded
kafka jobs-app	zookeeper jobs-app	forwarded
kafka jobs-app	zookeeper jobs-app	forwarded
coreapi jobs-app	elasticsearch jobs-app	forwarded

Cilium Policy Verdicts in Grafana



Using Hubble Observability Data for Security Compliance



Observability-Driven Zero Trust Approach



Day 1:

- identify commonly shared services to include in baseline policy.
- identify “low-hanging fruit” namespaces easily locked down at ingress/egress (no per-namespace exceptions required).
- identify the set of “exceptions” required for a given per-namespace policy

Day 2:

- Flow data can be sourced from dev/staging clusters as well as production to identify the network policy changes required to ship a new app version.
- Historical flow data can be analyzed to predict impact of new more stringent baseline or per-namespace policies.
- Historical flow data combined with enforced policies can be used to detect overly broad rules. For example, “egress allow 0.0.0.0/0 port 443” when a rule to a few specific DNS names would suffice.



eBPF resources



eCHO News is curated by Bill Mulligan

Bill Mulligan is working to grow the Cilium community

[FOLLOW ON TWITTER](#)



Join Cilium & eBPF on Slack.
10970 users are registered so far.

eCHO

eBPF YouTube podcast:

<https://www.youtube.com/channel/UCJFUxkVQTBJh3LD1wYBWvuQ>

eCHO News

Bi-weekly eBPF newsletter:

<https://cilium.io/newsletter/>

eBPF & Cilium Slack

<http://slack.cilium.io/>



ISOVALENT

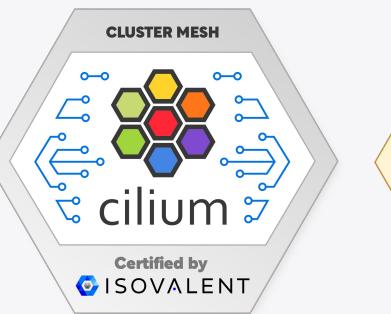
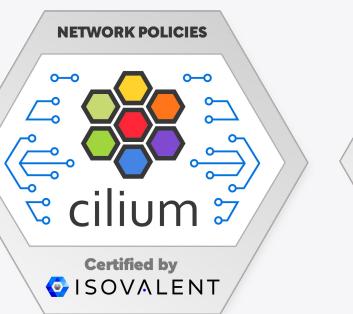


Practical Labs

... to become a Cilium & eBPF Jedi



isovalent.com/labs



Get badges 🏅





Zero Trust Visibility

Using Cilium, Hubble & Grafana for namespace security



<https://isogo.to/zero-trust-lab>



Workshops

Zurich — 1st June

Basel — 7th June

Geneva — 14th September

 isovalent.com/workshop-tour



ISOVALENT

Thank you!

